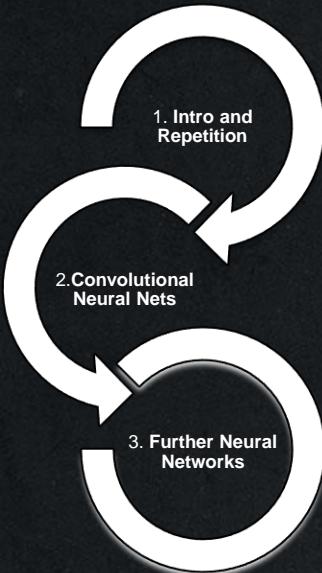


Scientific Machine and Deep Learning for Design and Construction in Civil Engineering

@ ETH Zürich 2023

Agenda



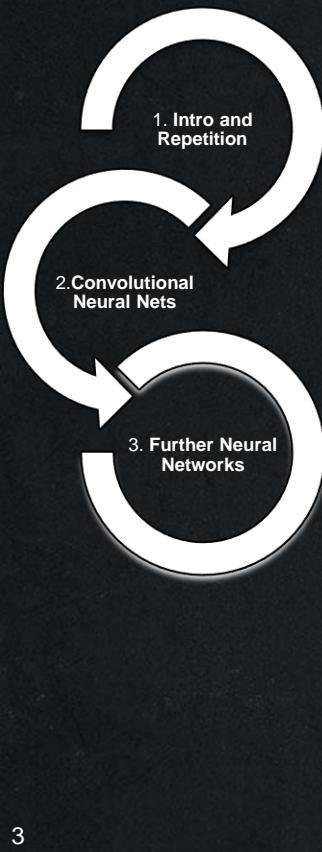
1 Intro and Repetition

2 Convolutional Neural Nets

3 Further Neural Networks

01. Introduction

Lectures / Project Consultation



Dr. Michael A. Kraus

- PhD with honors 2019
@ Bundeswehr University Munich
- Post-Doc @ Stanford University
- Senior-Research @ ETH Zürich
- Co-Lead Immersive Design Lab

Dr. Danielle Griego

- PhD 2020
@ ETH Zürich
- Post-Doc @ ETH Zürich
- Executive Director of Design++



Vera Balmer, M.Sc.

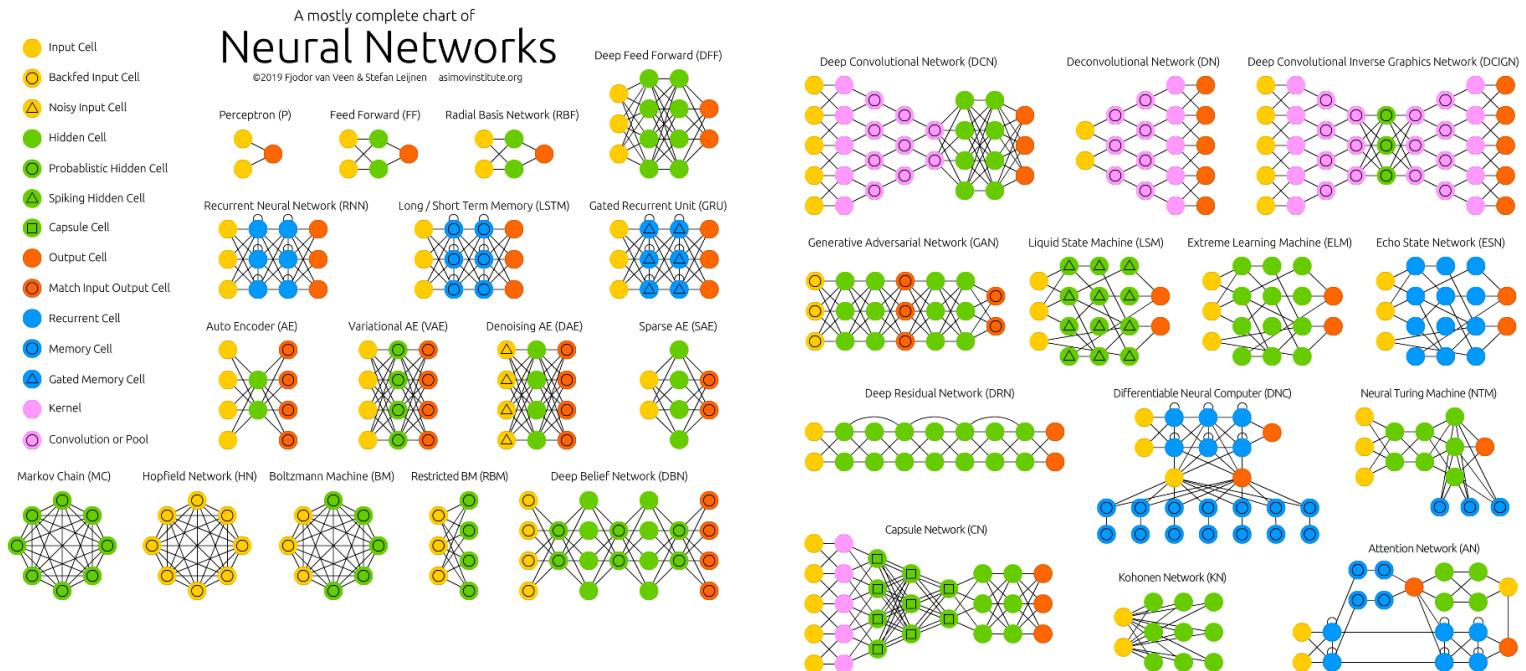
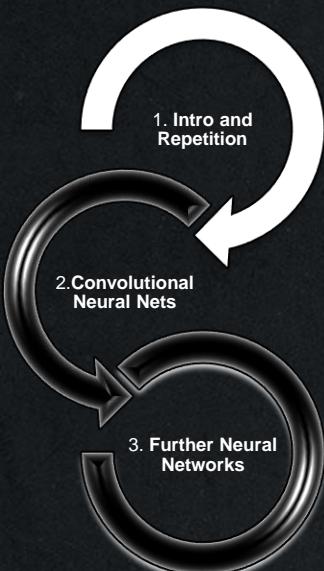
- M.Sc. Civil Engineering 2022
@ ETH Zürich
- PhD candidate @ ETH: IBK & AI Center



01. Intro and Repetition

Deep Learning Intro

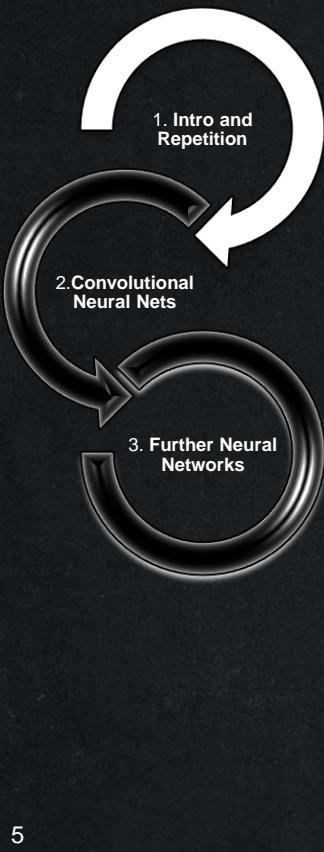
The neural network zoo



<https://www.asimovinstitute.org/neural-network-zoo/>

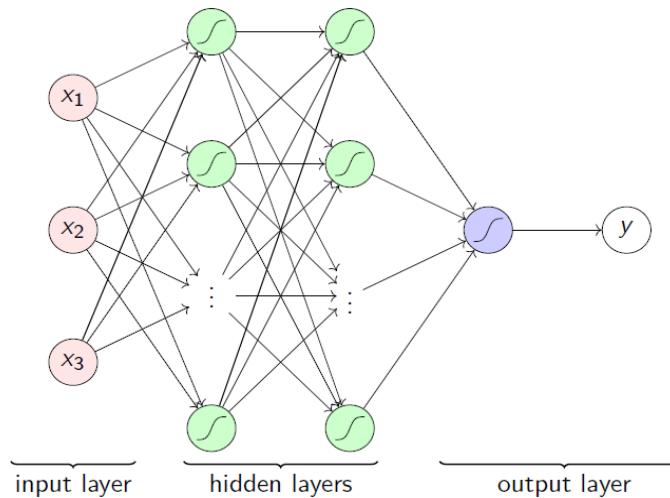
01. Intro and Repetition

Deep Learning Intro



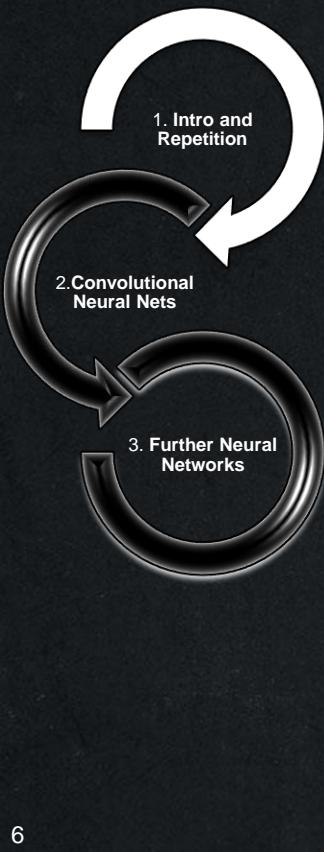
Repetition / Motivation

- So far: fully connected layers – each input is connected to each node
- Very powerful: can represent any kind of (linear) relationship between inputs



01. Intro and Repetition

Deep Learning Intro

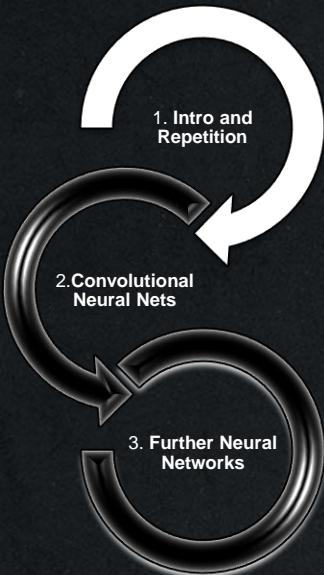


Repetition / Motivation

- So far: fully connected layers – each input is connected to each node
- Very powerful: can represent any kind of (linear) relationship between inputs
- Large part of ML: images / videos / sounds
- Assume we have:
 - an image with size 512×512 pixels
 - one hidden layer with 8 neurons
 - $(512^2 + 1) \times 8 > \text{2 million trainable weights}$

01. Intro and Repetition

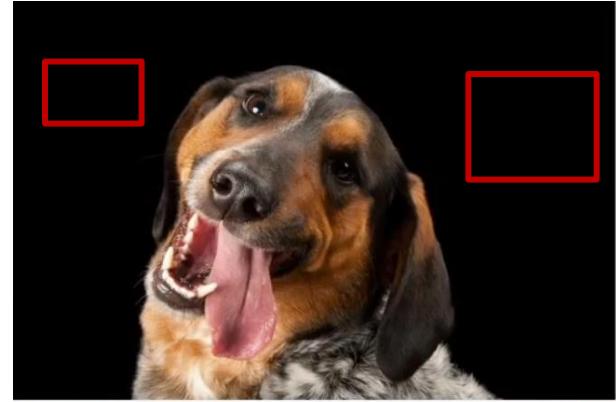
Deep Learning Intro



Repetition / Motivation

The size is a problem.
Is there something else?

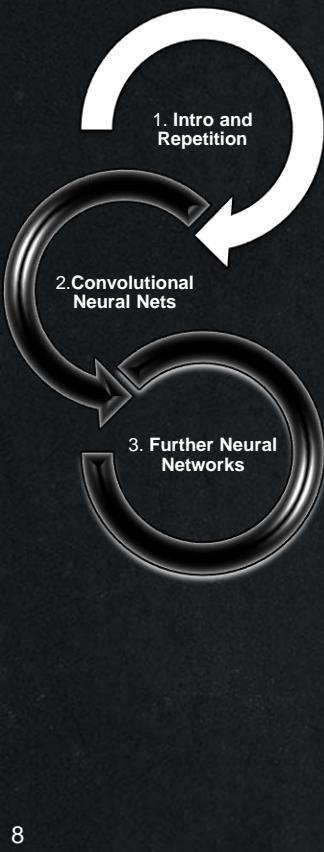
- Example: classify between dog and cat
- **Pixels are bad features**
 - Highly correlated
 - Scale dependent
 - Intensity variations
 -
- **Pixels are a bad representation from a ML point of view**



Source: <https://news.nationalgeographic.com>

01. Intro and Repetition

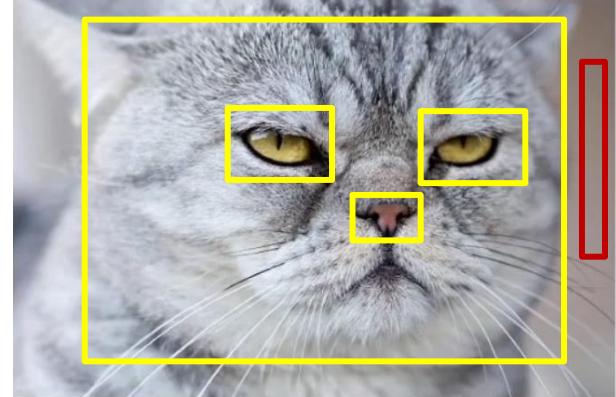
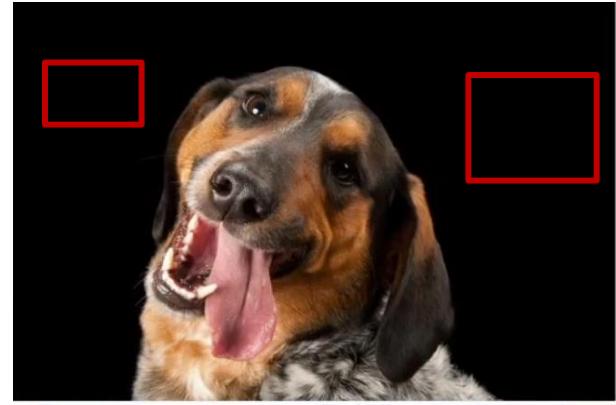
Deep Learning Intro



Repetition / Motivation

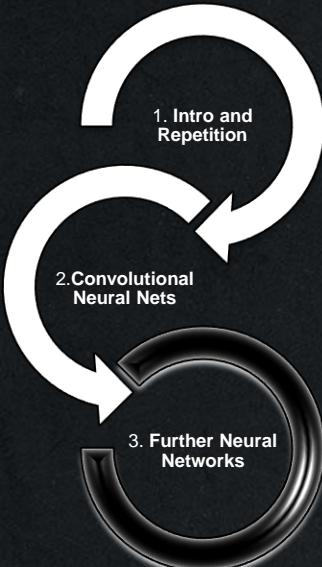
Can we find a better representation?

- We have a certain degree of locality in an image
- We can find the same „macro features“ at different locations
- Hierarchy of features:
 - Edges + corners => eyes
 - Eyes + nose + ears => face
 - Face + body + legs => animal
- **Composition matters!!!**
- **Learn better representation, then classify!**



Source: <https://news.nationalgeographic.com>

Agenda



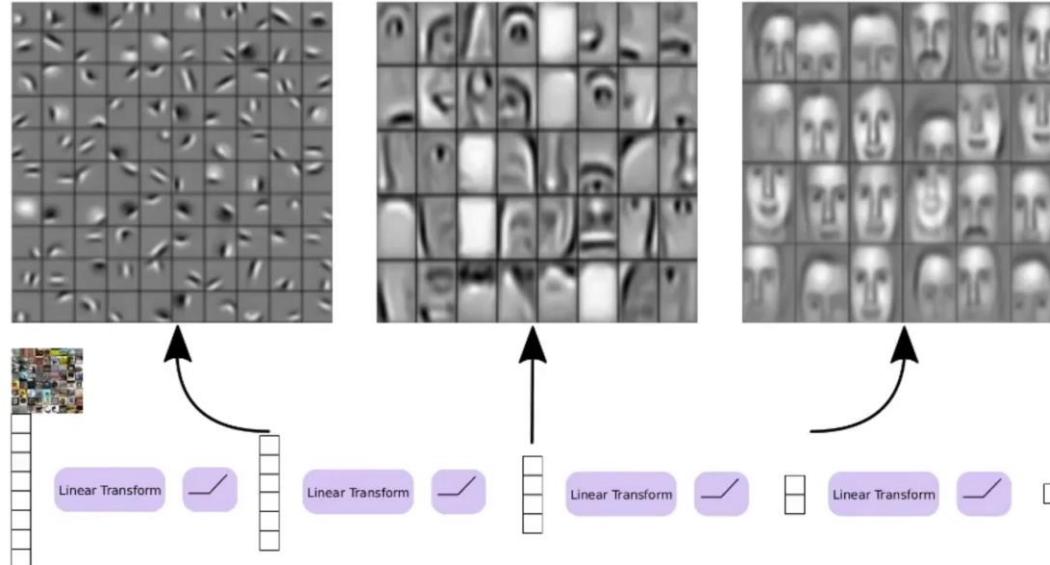
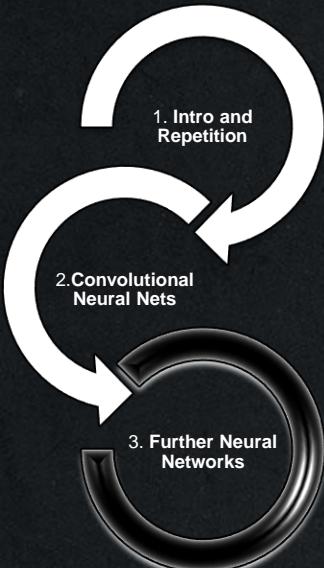
1 Intro and Repetition

2 Convolutional Neural Nets

3 Further Neural Networks

02. Convolutional Neural Networks

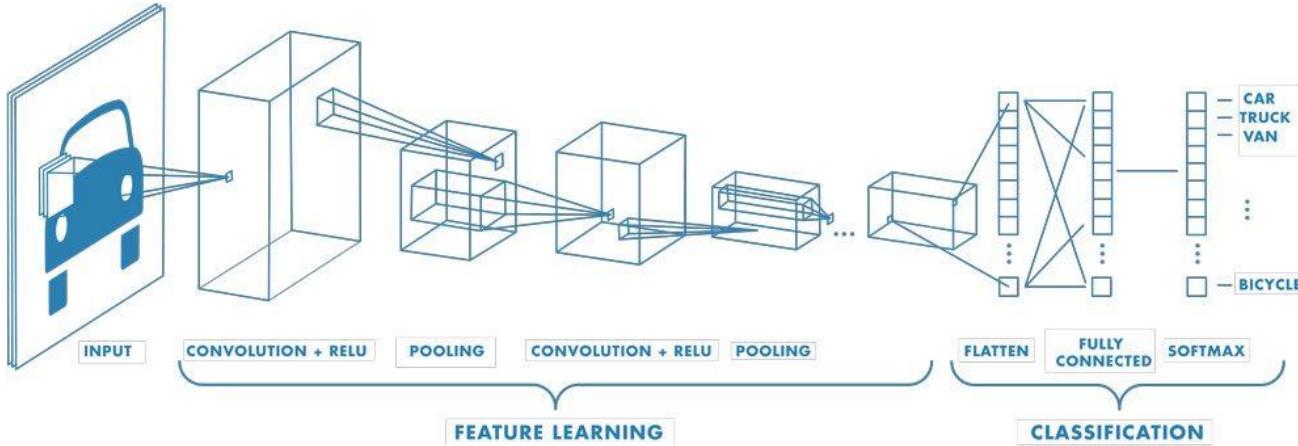
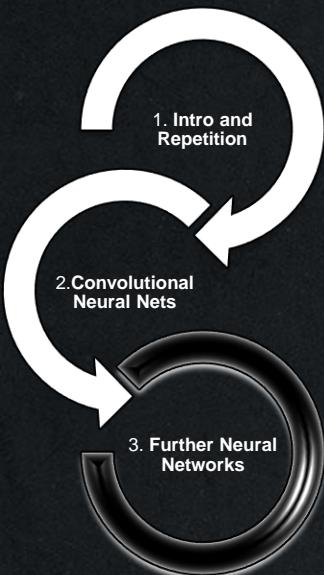
CNN



- **local** connectivity => filters
- use **same filters** over the whole image => translational equivariance
- Hierarchy of filters working on **different scales**
- + learning = **Convolutional Neural Networks (CNN)**

02. Convolutional Neural Networks

CNN - Architecture

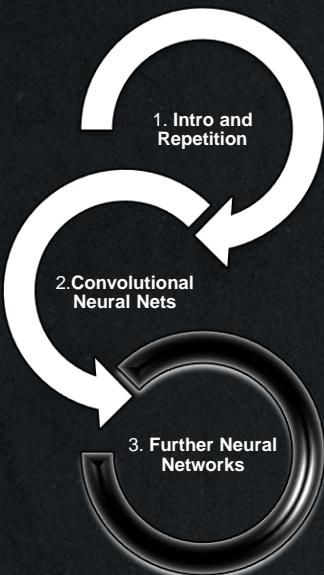


- Convolutional layer: feature extraction
- Activation function: nonlinearity
- Pooling layer: compress and aggregate information; save parameters
- Last layer: fully-connected for classification

<https://de.mathworks.com/discovery/convolutional-neural-network-matlab.html>

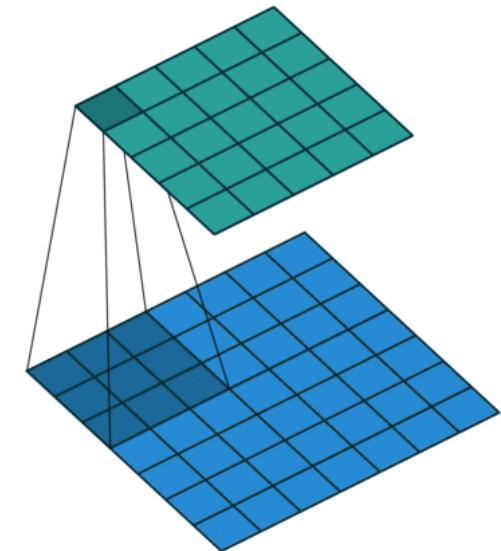
02. Convolutional Neural Networks

CNN - Anatomy



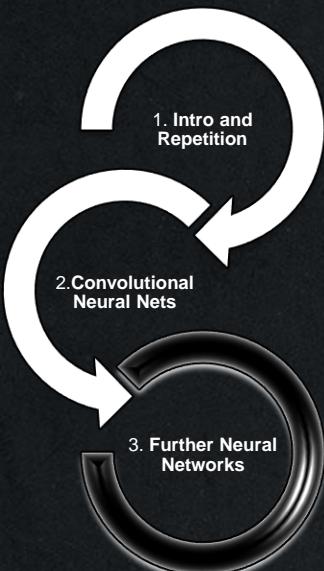
Convolutional Layer – Local Connectivity

- Exploits spatial structure by only connecting pixels in a neighbourhood
- can be expressed as fully connected layer: **except** for local connections, each entry in **W** is zero
- effective weights:
filter of size $3 \times 3, 5 \times 5, 7 \times 7, \dots$
- Features, that are important at one location are likely important anywhere in the image
- **convolution with trainable filters**



02. Convolutional Neural Networks

CNN - Anatomy



Convolution

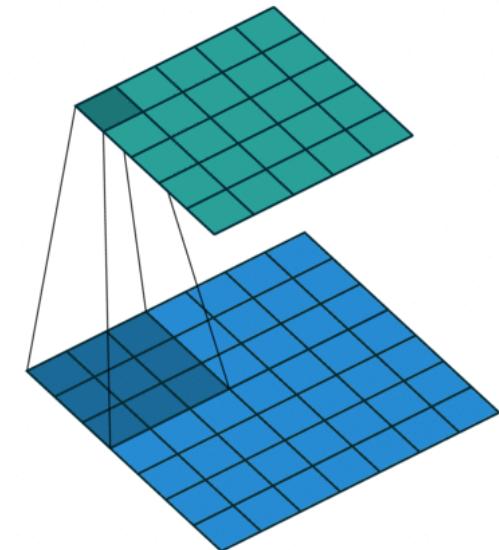
- Convolution:

$$(f * g)(x) = \int_{-\infty}^{\infty} f(\tau)g(x - \tau)d\tau$$

- Cross-correlation:

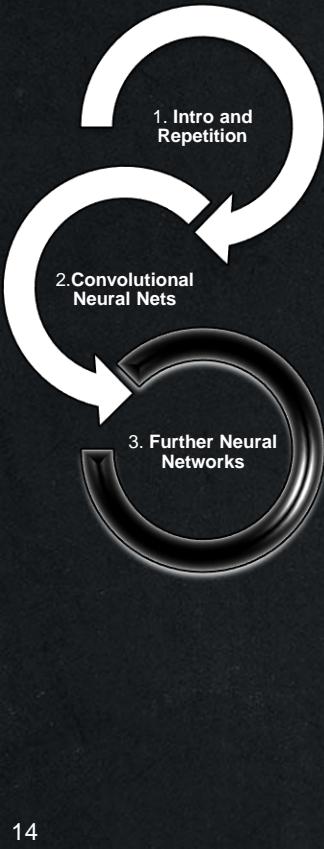
$$(f \times g)(x) = \int_{-\infty}^{\infty} f(\tau)g(x + \tau)d\tau$$

- Cross-correlation is convolution with a flipped kernel g – and vice versa
- Implementation: cross-correlation is frequently used in the forward pass – the weights are initialized randomly anyway



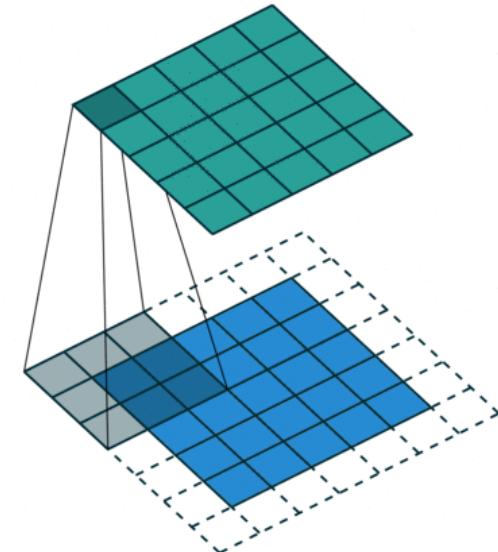
02. Convolutional Neural Networks

CNN - Anatomy



Padding

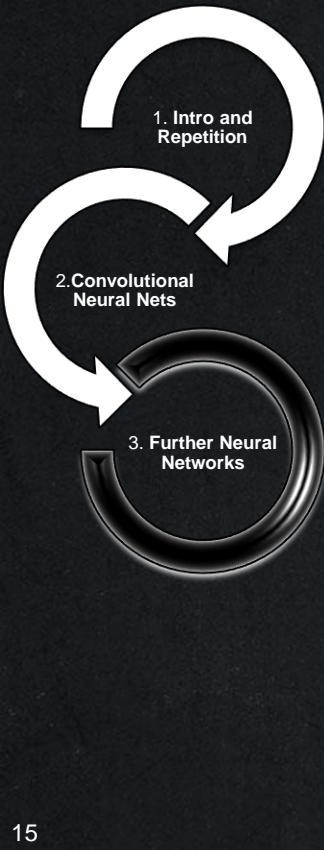
- Convolution reduces image size by $2 \left[\frac{n}{2} \right]$ pixels (n : kernel size)
- Necessary to pay attention to the borders:
- „same“ padding (usually „zero“ padding“):
=> Input and output have the same size
- „valid“ / no padding:
=> output is smaller than the input



https://github.com/vdumoulin/conv_arithmetic

02. Convolutional Neural Networks

CNN - Anatomy

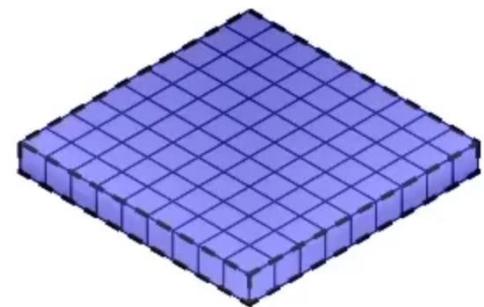
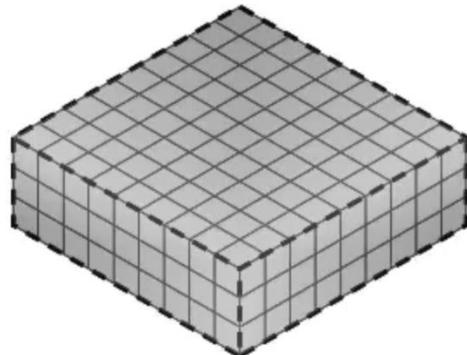


Forward Pass: Multi-channel convolutions

- Input of size $X \times Y \times S$
- H filters with size $M \times N \times S$

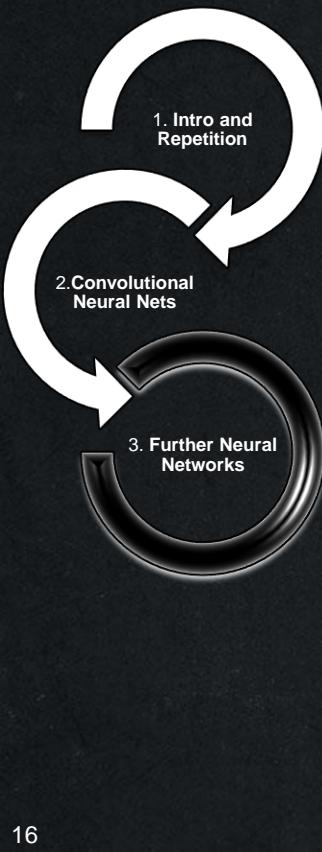
S : number of input channels

fully connected across channels



02. Convolutional Neural Networks

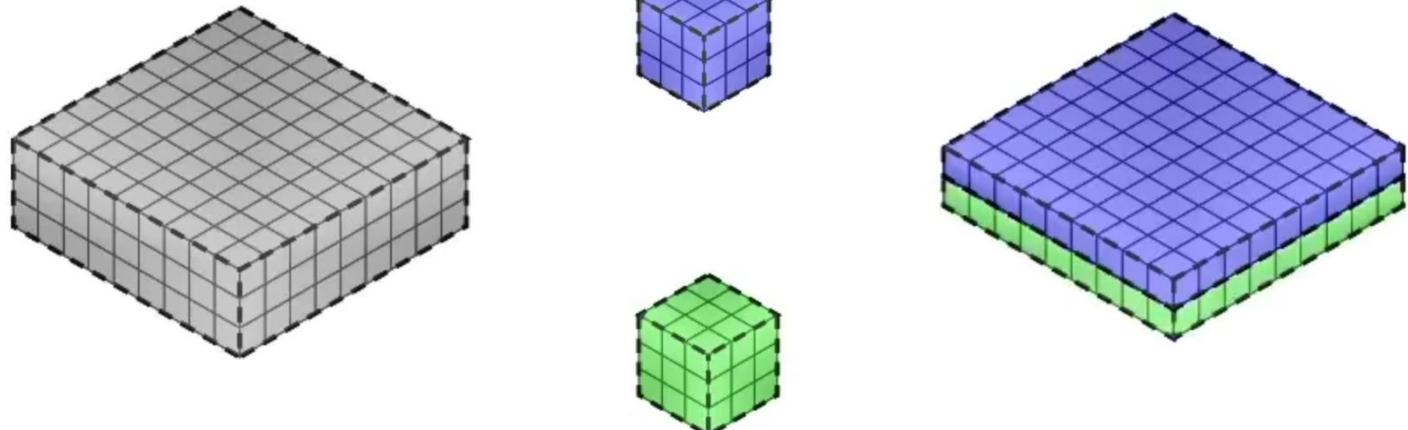
CNN - Anatomy



Forward Pass: Multi-channel convolutions

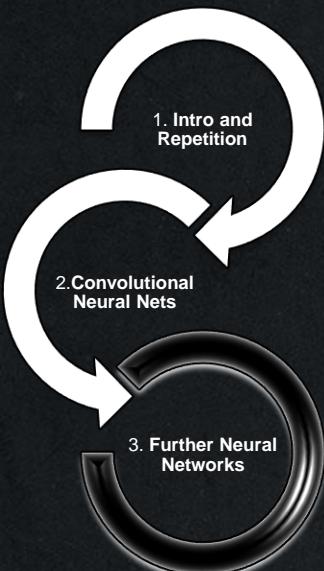
- Input of size $X \times Y \times S$
- H filters with size $M \times N \times S$
- Output dimensions: $X \times Y \times H$

S : number of input channels
fully connected across channels
with „same“ padding



02. Convolutional Neural Networks

CNN - Anatomy



Backward Pass: Multi-channel convolutions

- convolution expressed as matrix multiplication with matrix \mathbf{W} : using a Toeplitz matrix
- we can use the **same formulas** as for the fully connected layer

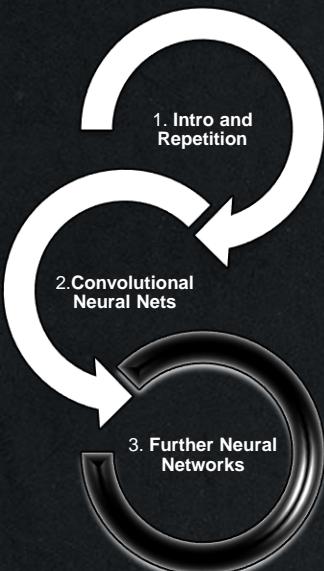
$$\begin{aligned} \mathbf{E}_{L-1} &= \mathbf{W}^T \mathbf{E}_L \\ \nabla \mathbf{W} &= \mathbf{E}_L \mathbf{X}^T \end{aligned}$$

where

- \mathbf{X} is the input
- \mathbf{E}_j is the error in layer $j \in \{L, L - 1\}$

02. Convolutional Neural Networks

CNN - Anatomy

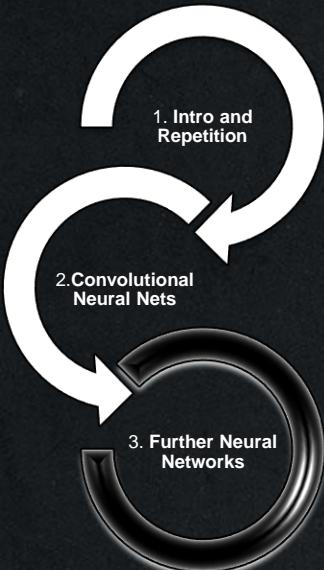


Convolutional Layer: What have we gained?

- stack multiple filters to get a trainable filter bank
- layer with 8 filters (nodes) with 5×5 neighborhood
 - $5^2 \times 8 = 200$ weights
- Convolution: independend of image size!
- Much more training data for one weight!

02. Convolutional Neural Networks

CNN - Anatomy

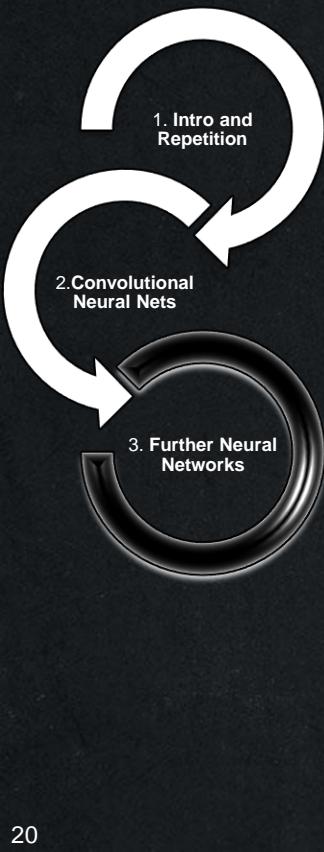


Convolutional Layer: Strided Convolution

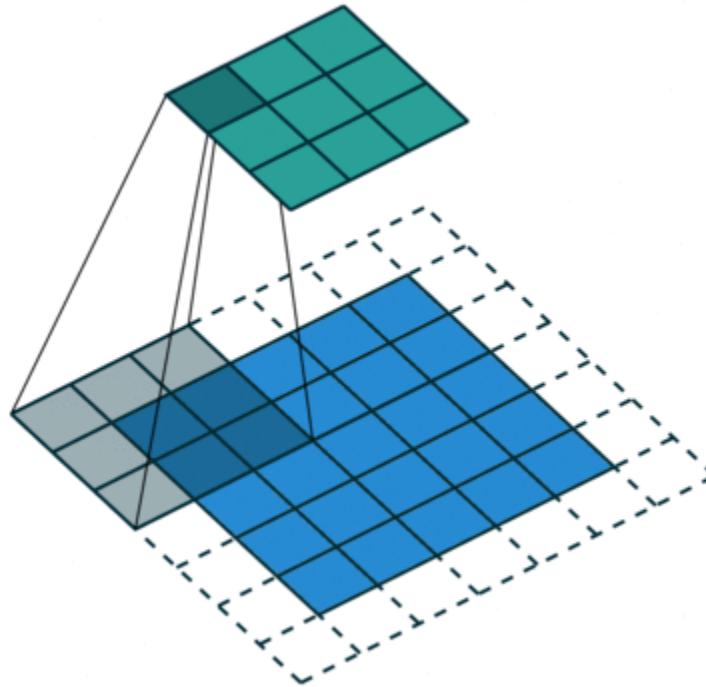
- instead of multiplying the filter at each pixel position, we can skip some positions
- stride s describes the offset
- reduces the size of the output by a factor of s
- *Mathematically:* *convolution and subsampling at the same time*

02. Convolutional Neural Networks

CNN - Anatomy



Convolutional Layer: Strided Convolution

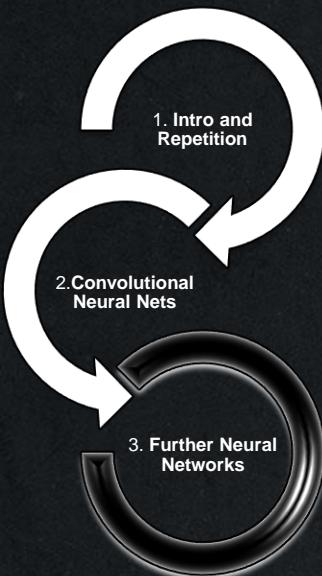


https://github.com/vdumoulin/conv_arithmetic

strided convolution with stride $s = 2$

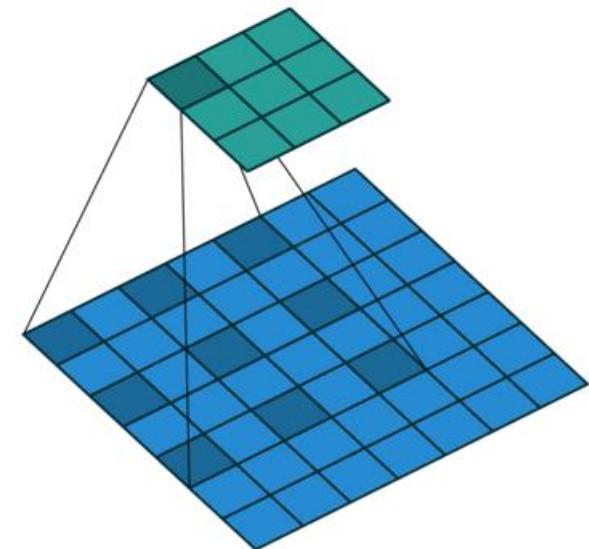
02. Convolutional Neural Networks

CNN - Anatomy



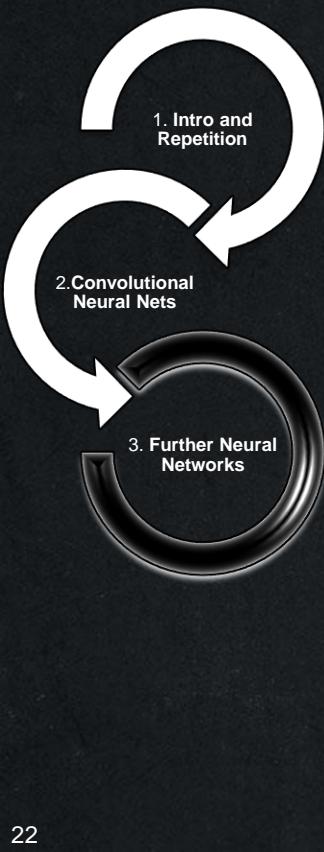
Convolutional Layer: Dilated / Atrous Convolution

- Additional variant of convolution in neural networks
- Dilate convolution kernel: skip certain pixels
- Goal: wider receptive field with less parameters / weights



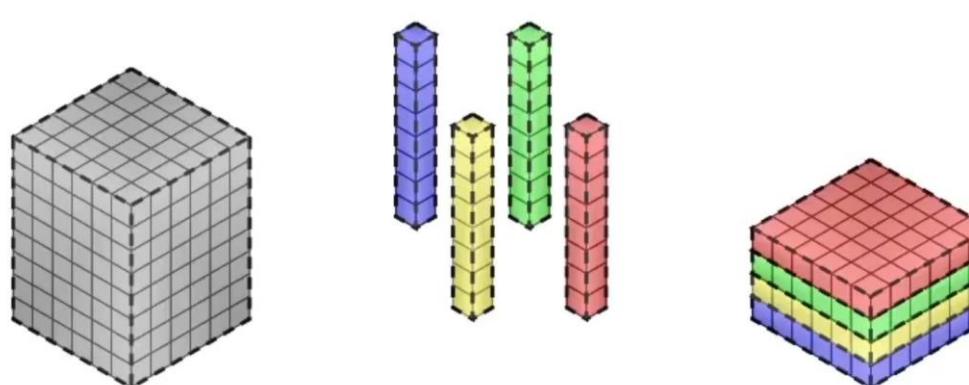
02. Convolutional Neural Networks

CNN - Anatomy



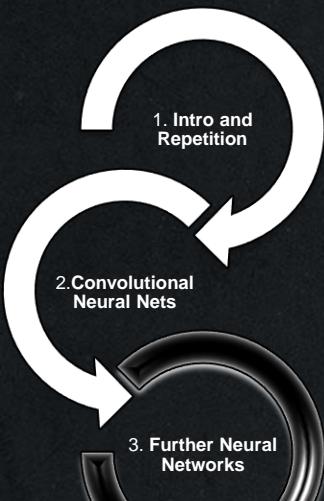
Convolutional Layer: 1 x 1 convolution concept

- So far: H filters with neighbourhood 3×3 , 5×5 , ... and depth S
- Filters are fully connected in „depth“ direction
- We can decrease the neighborhood to 1×1
- And just use the fully connected property in depth dimension

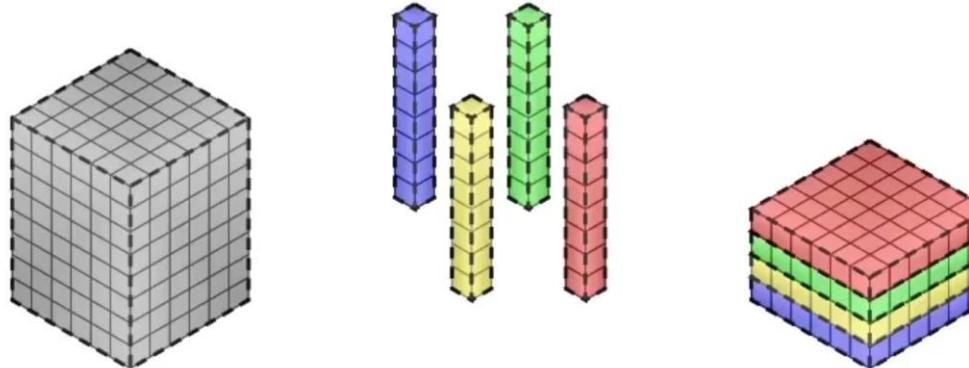


02. Convolutional Neural Networks

CNN - Anatomy



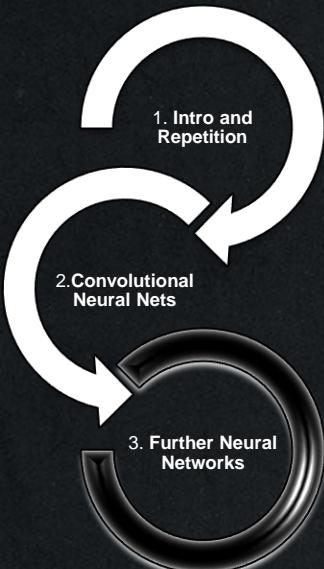
Convolutional Layer: 1×1 convolution concept



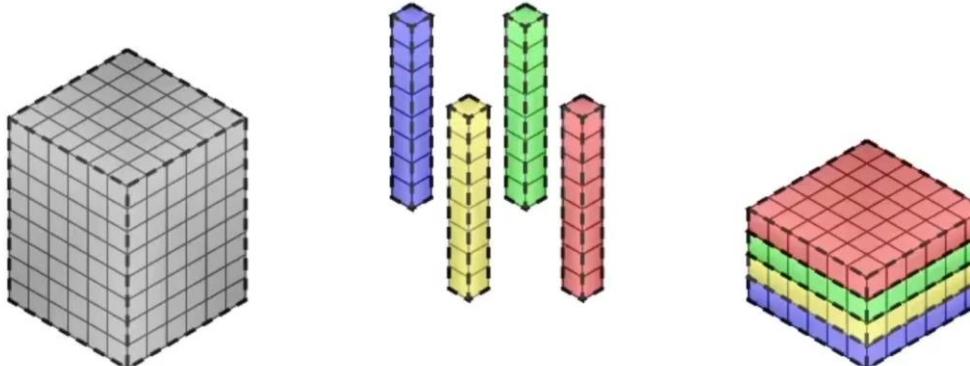
- Dimension reduction / expansion from S channels to H channels
- **If we flatten the input, 1×1 convolutions are fully connected layers**

02. Convolutional Neural Networks

CNN - Anatomy



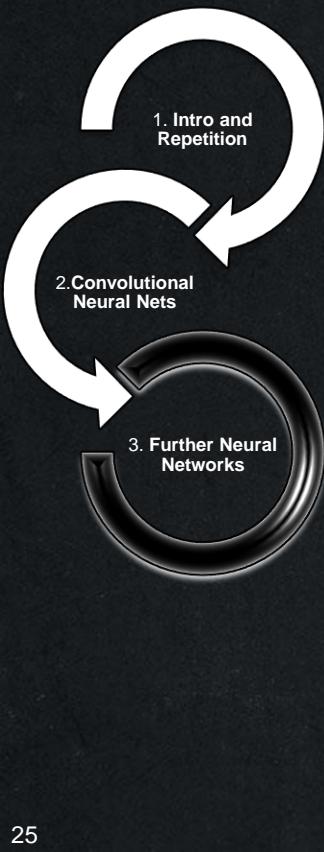
Convolutional Layer: 1×1 convolution concept



- Dimension reduction / expansion from S channels to H channels
- **if we flatten the input, 1×1 convolutions are fully connected layers**
- Simple and effective method to decrease the size of a network
- **Learns** dimensionality reduction, e.g. reduce redundancy in your feature maps

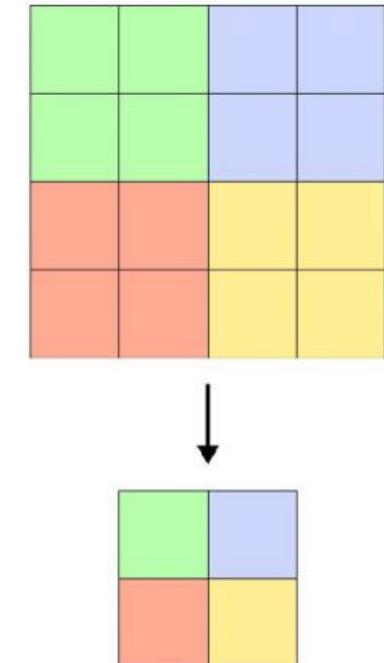
02. Convolutional Neural Networks

CNN - Anatomy



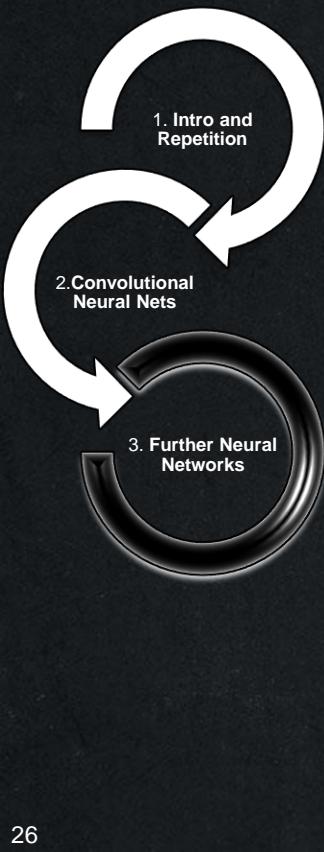
Convolutional Layer: Idea behind pooling layers

- fuses information of input across spacial locations
- decreases number of parameters
- reduces computational costs and overfitting
- assumptions:
 - features are hierarchically structured
 - creates „summaries“ of regions
 - provides translational invariance
 - exact location of a feature is not important



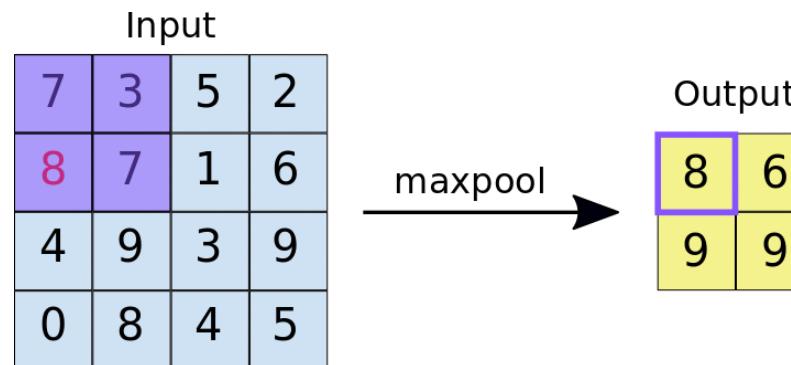
02. Convolutional Neural Networks

CNN - Anatomy



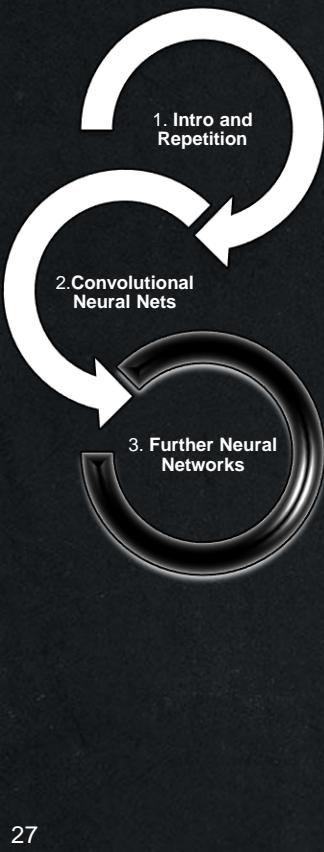
Convolutional Layer: Max-Pooling – Forward Pass

- Propagate maximum value in a neighborhood to next layer
- Typical choices: 2×2 or 3×3 neighborhood
- „stride“ of pooling usually equals the neighborhood size
- Maximum propagation adds additional non-linearity



02. Convolutional Neural Networks

CNN - Anatomy



Convolutional Layer: Average-Pooling – Forward Pass

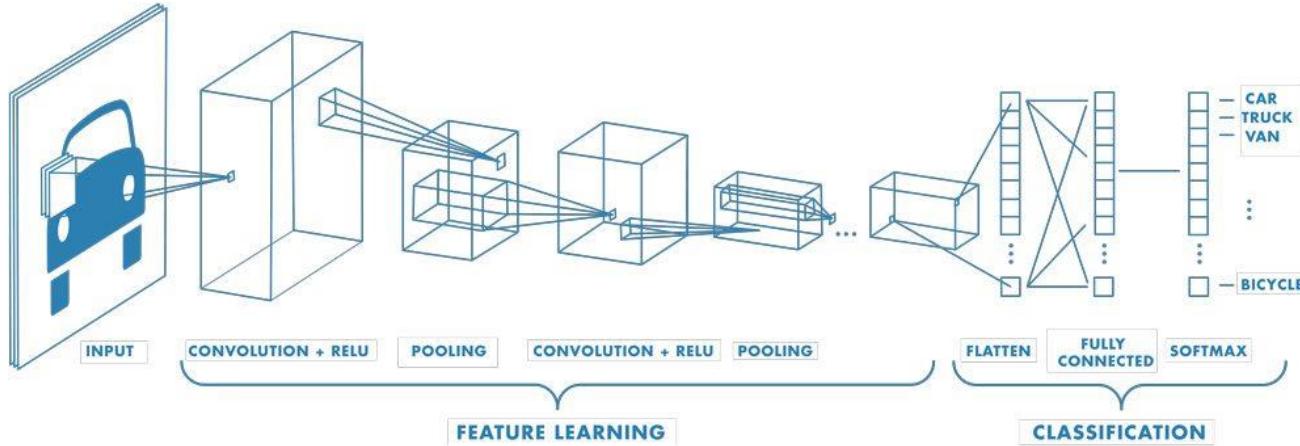
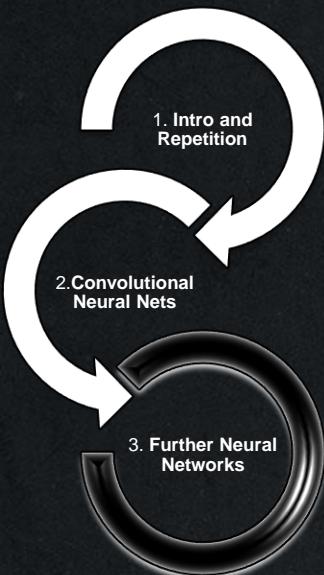
- Propagate average of a neighborhood
- Does not consistently perform better than max-pooling
- Backward pass: error is shared to equal parts

Convolutional Layer: further Pooling strategies

- Propagate average of a neighborhood
- Fractional max pooling
- Stochastic pooling
-

02. Convolutional Neural Networks

CNN - Architecture



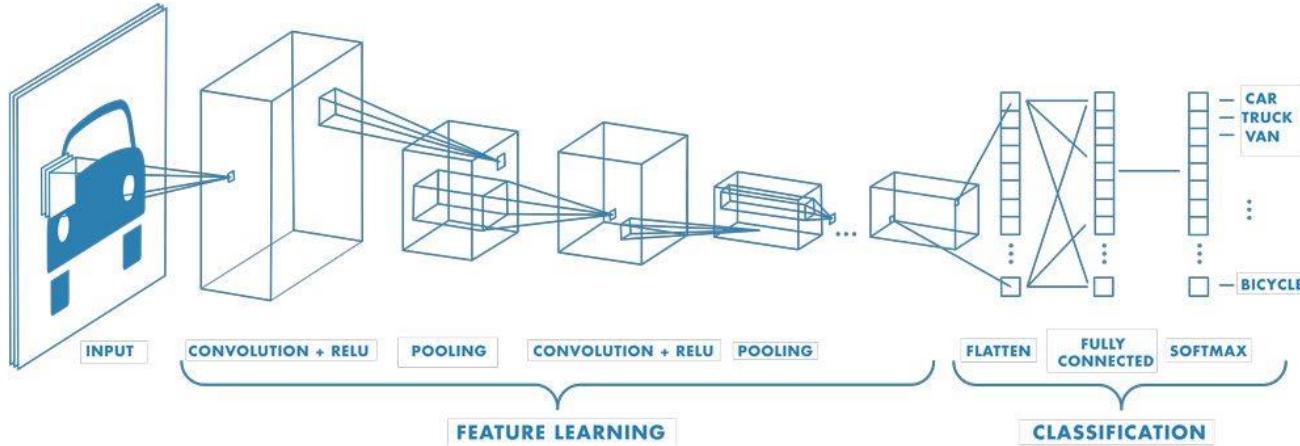
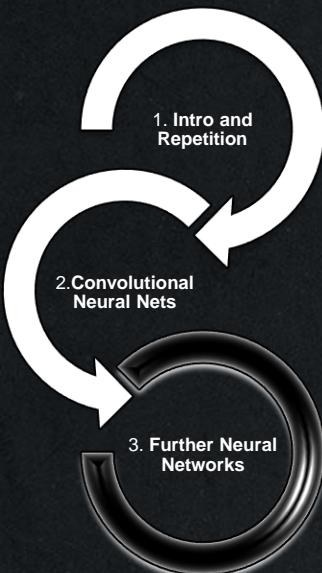
Four essential building blocks:

- Convolutional layer: feature extraction
- Activation function: nonlinearity
- Pooling layer: compress and aggregate information; save parameters
- Last layer: fully-connected for classification

<https://de.mathworks.com/discovery/convolutional-neural-network-matlab.html>

02. Convolutional Neural Networks

CNN - Architecture



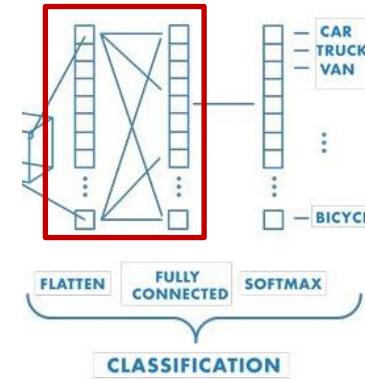
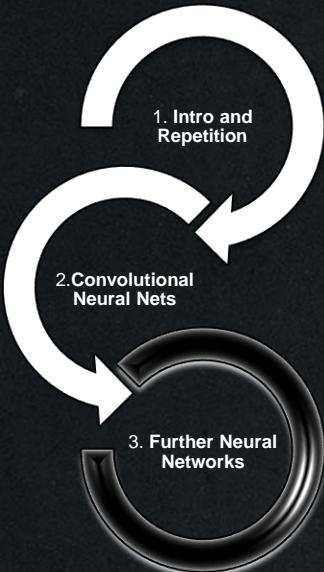
Four Three essential building blocks:

- Convolutional layer: feature extraction
- Activation function: nonlinearity
- Pooling layer: compress and aggregate information; save parameters
- Last layer fully-connected for classification ***can be replaced***

<https://de.mathworks.com/discovery/convolutional-neural-network-matlab.html>

02. Convolutional Neural Networks

CNN - Architecture

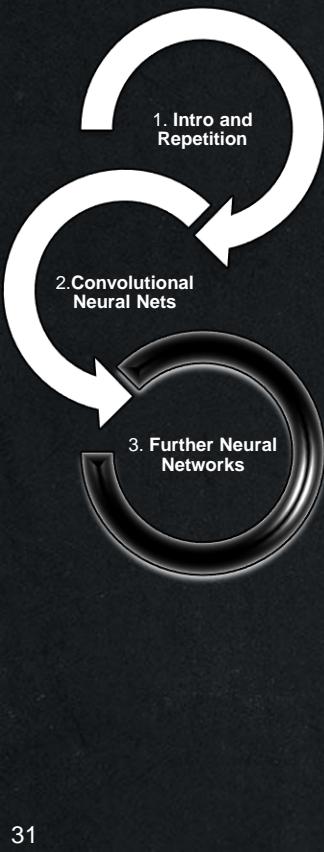


<https://de.mathworks.com/discovery/convolutional-network-matlab.html>

- convolutional and pooling layers generate better representations => better features
- fully connected layers for classification
- **alternatively and equivalently:** use flatten & 1×1 convolution or $N \times N$ convolution
- enables **arbitrary input sizes** in combination with global average pooling

02. Convolutional Neural Networks

Regularization



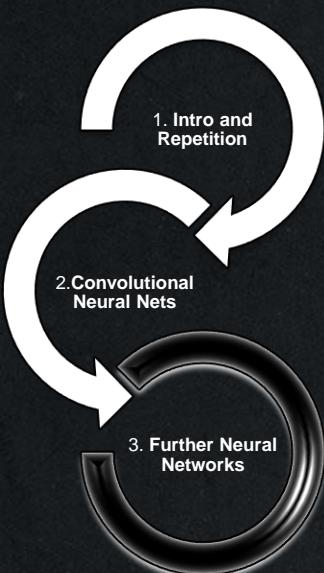
Regularization

	Underfitting	Just right	Overfitting
Symptoms	<ul style="list-style-type: none">• High training error• Training error close to test error• High bias	<ul style="list-style-type: none">• Training error slightly lower than test error	<ul style="list-style-type: none">• Very low training error• Training error much lower than test error• High variance
Regression illustration			
Classification illustration			
Deep learning illustration			
Possible remedies	<ul style="list-style-type: none">• Complexify model• Add more features• Train longer		<ul style="list-style-type: none">• Perform regularization• Get more data

<https://stanford.edu/~shervine/teaching/cs-229/cheatsheet-machine-learning-tips-and-tricks>

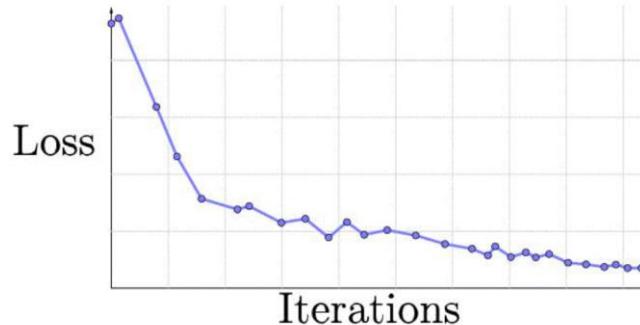
02. Convolutional Neural Networks

Regularization

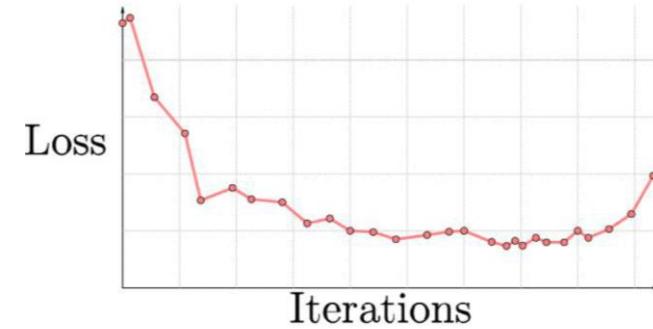


Regularization – Overfitting on Learning Curve

Training set



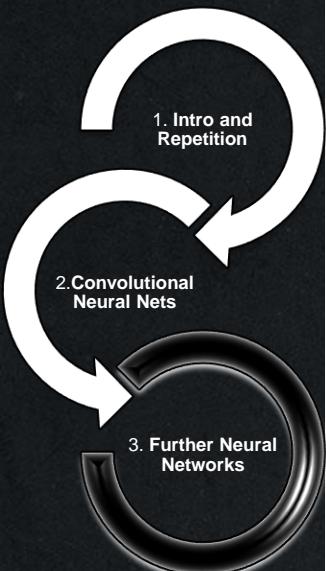
Test set



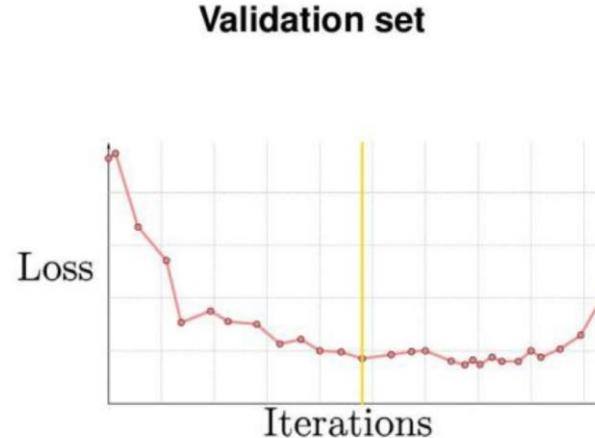
- test set **must not** be used for training!
=> split of **validation set** from training data

02. Convolutional Neural Networks

Regularization



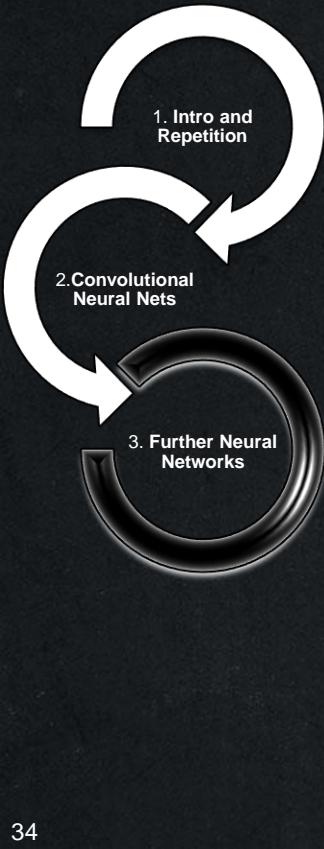
Regularization – Early Stopping



- Define stopping criterion
 - => use parameters with **minimum validation loss**

02. Convolutional Neural Networks

Regularization



Regularization – Data Augmentation

=> Artificially enlarge dataset

- but how?
- Every transformation which the label should be invariant to:

Probably the same class

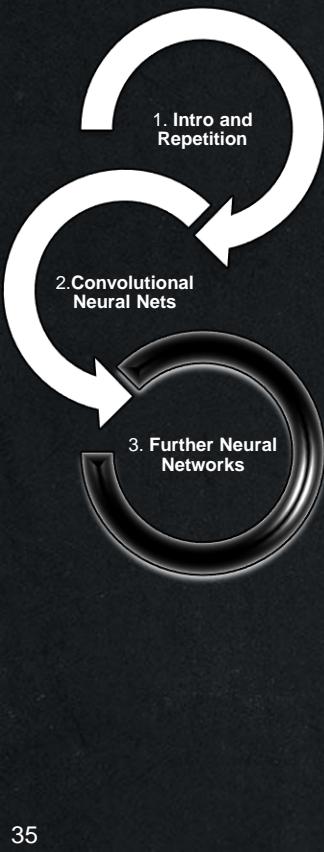


still you have to be careful



02. Convolutional Neural Networks

Regularization

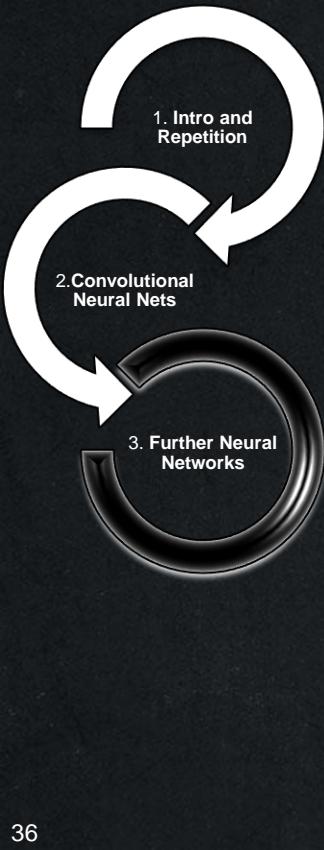


Regularization – Data Augmentation: Common transformations

- Random spatial transformations:
 - affine transformation
 - elastic transformation
 - ...
- Pixel transformation
 - changing resolutions
 - Random noise
 - Changing pixel distribution
 -

02. Convolutional Neural Networks

Regularization



Regularization – Maximum a Posteriori Estimate

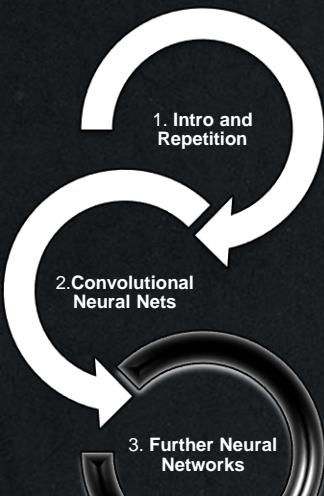
- Bayesian approach considering the weights uncertainty: $\mathbf{w} \mapsto p(\mathbf{w})$
- and we have a dataset \mathbf{X} with associated labels \mathbf{Y}
- We obtain the Maximum a Posteriori (MAP) estimate:

$$MAP(\mathbf{w}) = \max_{\mathbf{w}} \{p(\mathbf{Y}|\mathbf{X}, \mathbf{w})\} p(\mathbf{w})$$

MLE estimator prior

02. Convolutional Neural Networks

Regularization



Regularization – Maximum a Posteriori Estimate – enforce a small norm

$$\tilde{L}(\mathbf{w}, \mathbf{X}, \mathbf{Y}) = L(\mathbf{w}, \mathbf{X}, \mathbf{Y}) + \lambda \|\mathbf{w}\|_2^2$$

- Where $\lambda > 0$ we identify this as the Lagrangian function of

$$\tilde{L}(\mathbf{w}, \mathbf{X}, \mathbf{Y}) = \min_{\mathbf{w}} \{L(\mathbf{w}, \mathbf{X}, \mathbf{Y}) \quad \text{s.t.} \|\mathbf{w}\|_2^2 < \alpha\}$$

- with an unknown data-dependent α

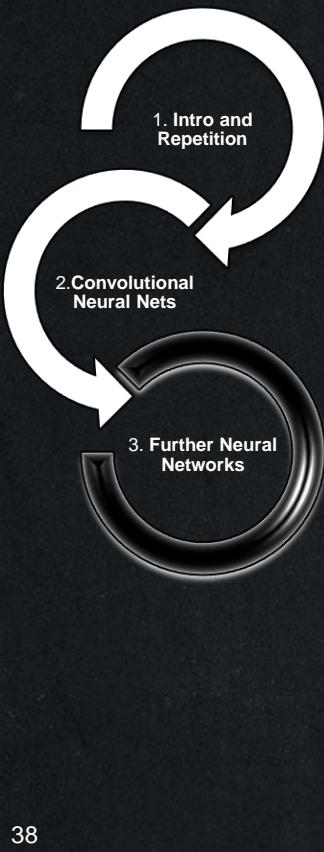
Backpropagation of the augmented loss

$$\mathbf{w}^{k+1} = (1 - \eta \lambda) \mathbf{w}^k - \eta \underbrace{\frac{\partial L}{\partial \mathbf{w}^k}}_{\text{shrinkage}}$$

- often called „weight decay“
- when optimizing the training loss for \mathbf{w} we will usually receive

02. Convolutional Neural Networks

Regularization

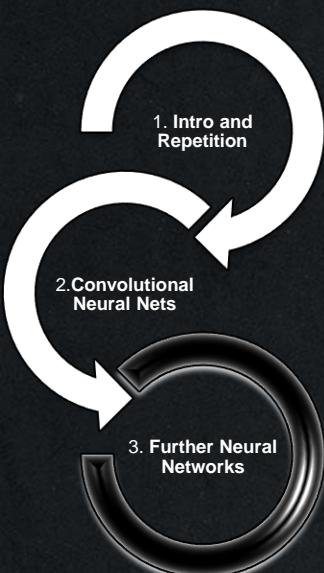


Regularization – other variants of changing the loss

- We can have a constraint and a λ individual for each layer
=> possible but no gains are reported
- instead of the weights, **activations** can be constrained
- different variants have been used for **sparse autoencoders**

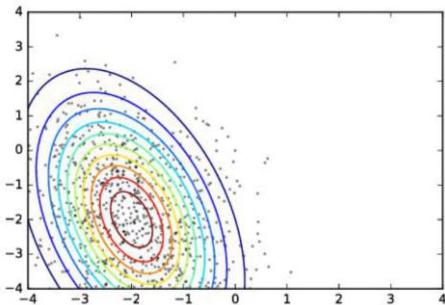
02. Convolutional Neural Networks

Regularization

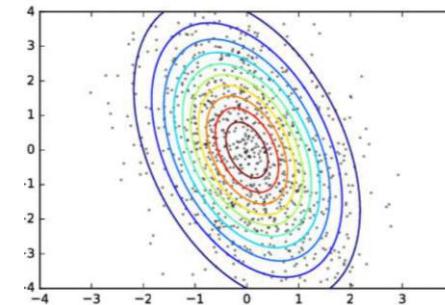


Data - Normalization

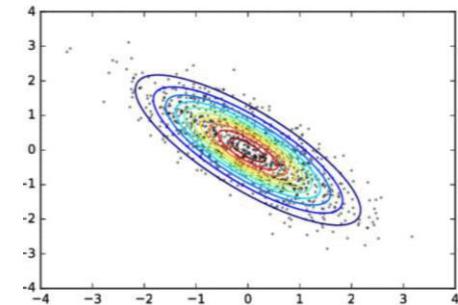
Original Data



mean substracted



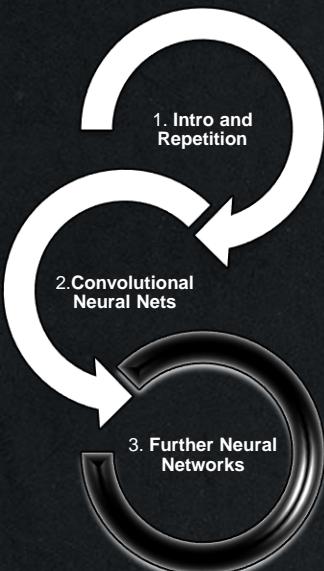
normalized variance



- Normalization: \min / \max or variance
- only use **training data** to compute normalization
- Normalization of input data
- Normalization **within** network

02. Convolutional Neural Networks

Regularization



Batch - Normalization

- Normalization as a new layer with 2 parameters and

$$\mathbf{x}_i = \frac{\mathbf{x}_i - \mu_{B,i}}{\sqrt{\sigma_{B,i}^2 + \epsilon}}$$

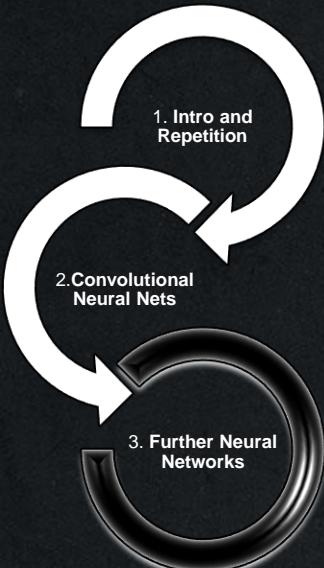
- and from **mini-batch**

$$\hat{y}_i = \gamma_i \mathbf{x}_i + \beta_i$$

- test time: replace μ_B and σ_B with μ and σ of the **training set**
- Paired with **convolutional layers** batch normalisation is different by computing a scalar μ and σ for every channel

02. Convolutional Neural Networks

Regularization



Batch – Normalization – why does it help?

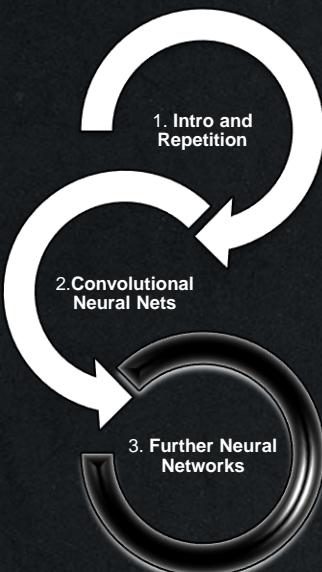
- overwhelming *empirical evidence* that batch normalization helps
- BN reduces internal covariate shift

Internal covariate shift:

- ReLu is not **zero-centered**
- Initialization and input distribution might not be normalized
 - => Input distributions **shift** over time
- **Deeper nets:** amplified effect
- Layers constantly adapt
 - => **Slow learning**

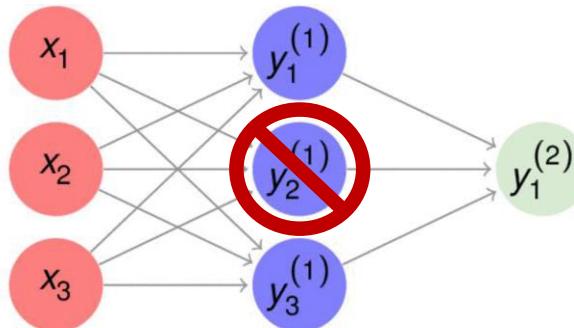
02. Convolutional Neural Networks

Regularization

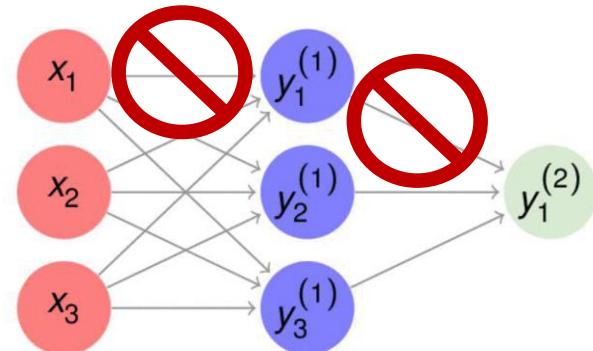


Dropout and Dropconnect

Dropout



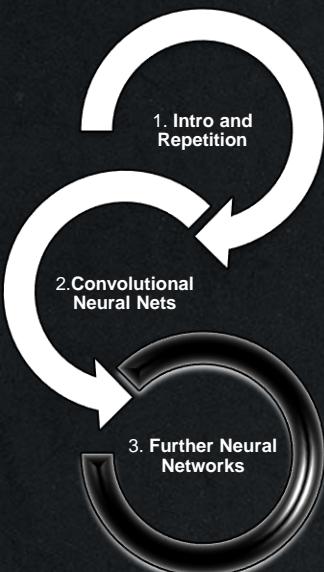
Dropconnect



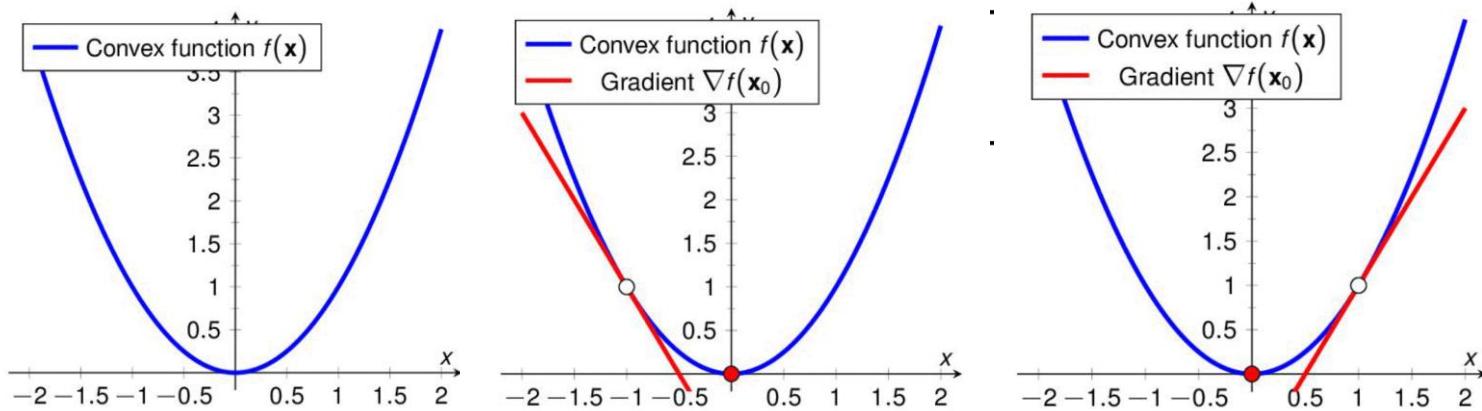
- Randomly set activations to zero with probability $1 - p$
- Test time: multiply activation with p
- Generalizes Dropout
- Less efficient implementation (masking)

02. Convolutional Neural Networks

Initialization



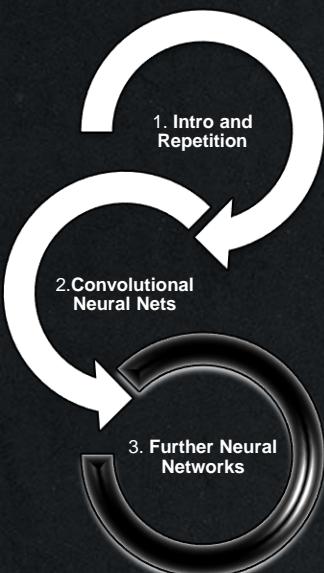
Initialization



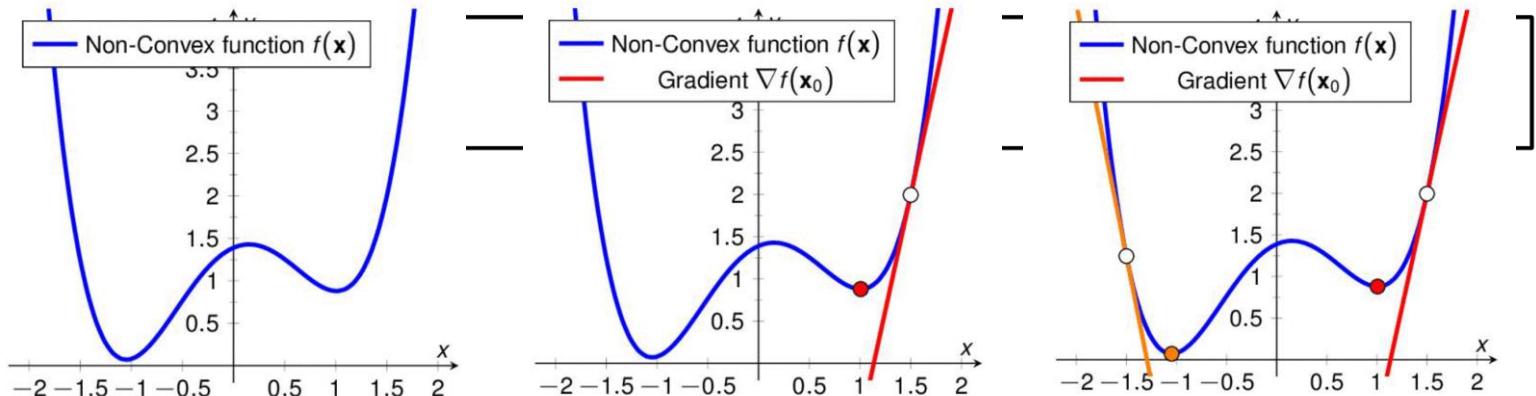
- Initialization **does not matter** in convex optimization problems

02. Convolutional Neural Networks

Initialization



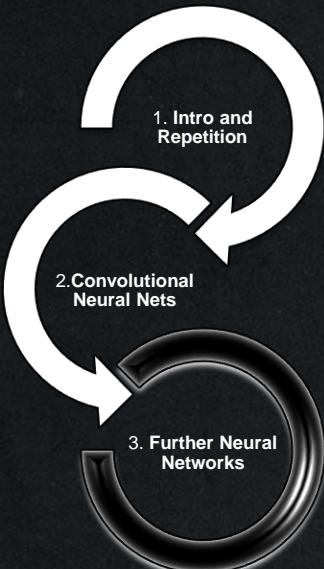
Initialization



- Initialization **does matter** in **non-convex** optimization problems
- Neural networks with a non-linearity in general lead to non-convex losses

02. Convolutional Neural Networks

Initialization



Initialization

Bias

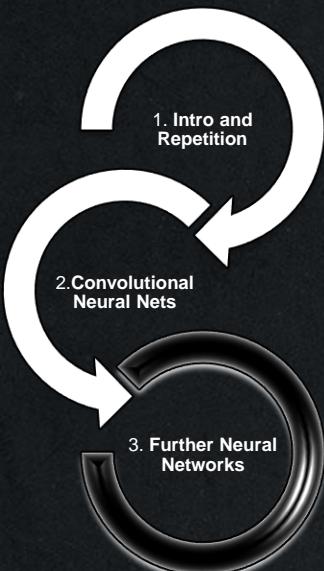
- Bias units can simply be **initialized to zero**
- Using **ReLU**: put a small positive constant (0.1) (=> dying ReLU issue)

Weights

- **weights** need to be **random** to **break symmetry**
- **!!! Very bad to initialize them with zeros (gradient then is zero)**
- Similar to learning rate their **variance** influences **stability** of learning
- **small** uniform / Gaussian values work

02. Convolutional Neural Networks

Initialization



Xavier Initialization

- We can calibrate the variances for the forward pass by initializing with a zero-mean Gaussian: $N(0, \sigma)$;

$$\sigma^2 = \frac{1}{fan_{in}} ; fan_{in} \text{ input dimensions of the weights}$$

- However the backward-pass:

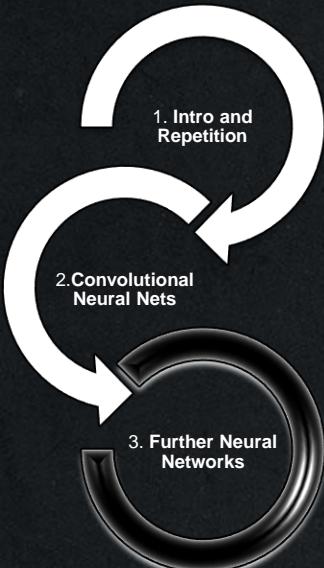
$$\sigma^2 = \frac{1}{fan_{out}} ; fan_{out} ; \text{ output dimensions of the weights}$$

- average those two:

$$\sigma = \sqrt{\frac{2}{fan_{out} + fan_{in}}}$$

02. Convolutional Neural Networks

Initialization



He Initialization and Conventional Initial Choices

- He initialization: assuming linear neurons is a problem

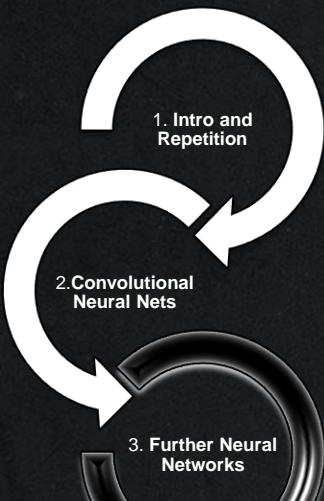
$$\text{for ReLu better use } \sigma = \sqrt{\frac{2}{fan_{in}}}$$

Conventional Initial Choices

- regularization
- Dropout using $p = 0.5$ for FCN; use selectively for CNN
- Mean subtraction
- Batch normalisation
- He initialisation

02. Convolutional Neural Networks

Transfer Learning

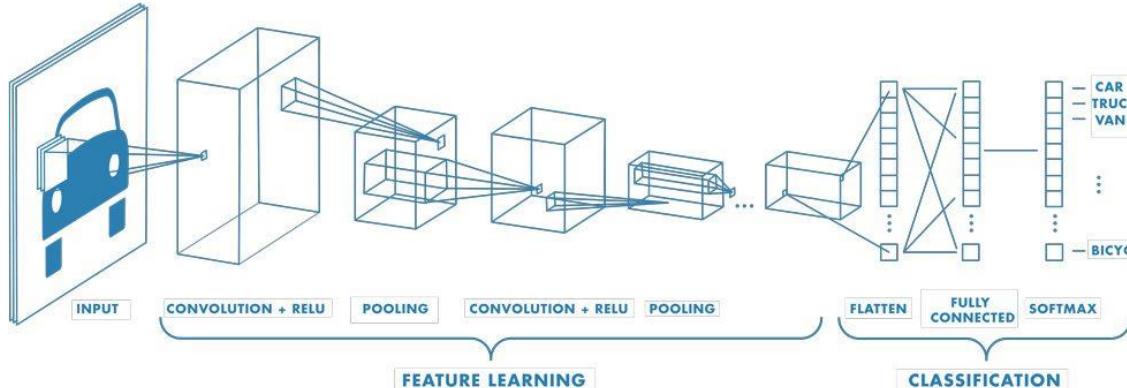
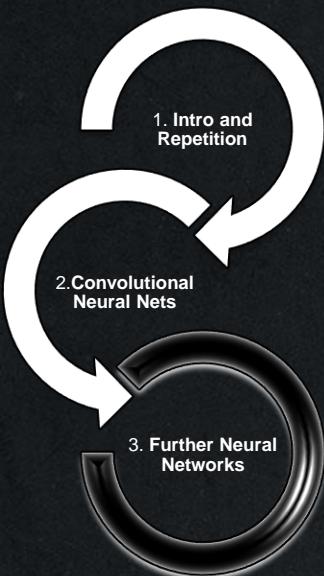


Transfer Learning

- Most technical / medical datasets are prohibitively small
- Reuse models trained on e.g. ImageNet
 - for a **different task** on the **same data**
 - for **different data** for the **same task**
 - for **different data** for a **different task**

02. Convolutional Neural Networks

Transfer Learning



<https://de.mathworks.com/discovery/convolutional-neural-network-matlab.html>

What to transfer?

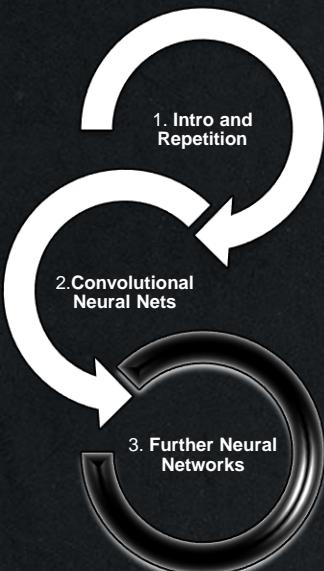
=> convolutional layers: feature extraction

=> expectation: less task-specific in earlier layers

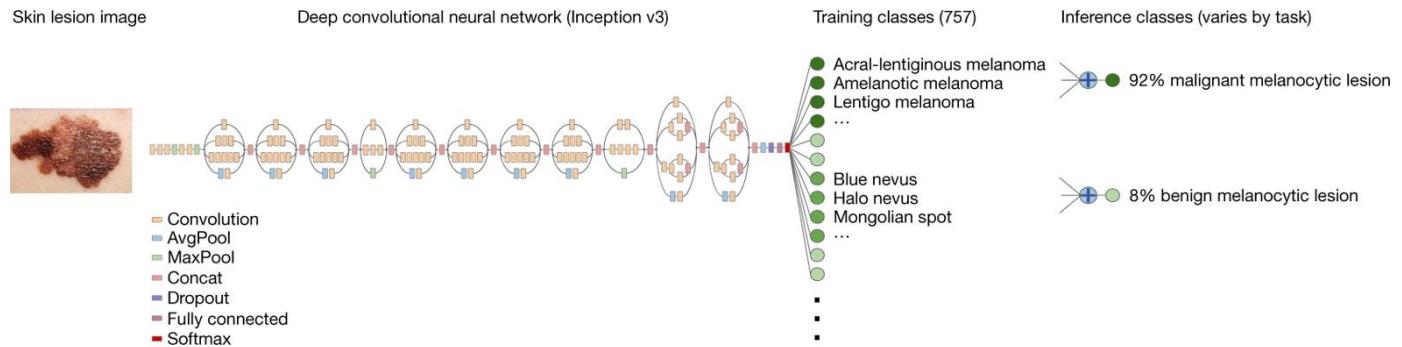
- We cut the network at some depth in the feature extraction part, extracted parts can be
 - fixed by setting $\eta = 0$
 - **fine-tuned**

02. Convolutional Neural Networks

Transfer Learning

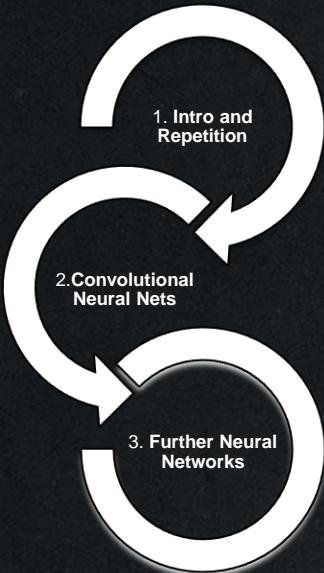


Example: Skin cancer detection



- State of the art architecture, pre-trained on ImageNet
- Fine-tuned on skin cancer data

Agenda



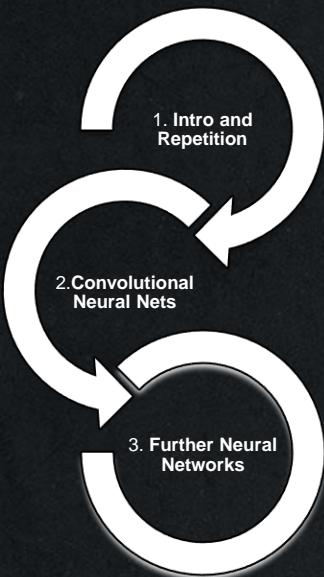
1 Intro and Repetition

2 Convolutional Neural Nets

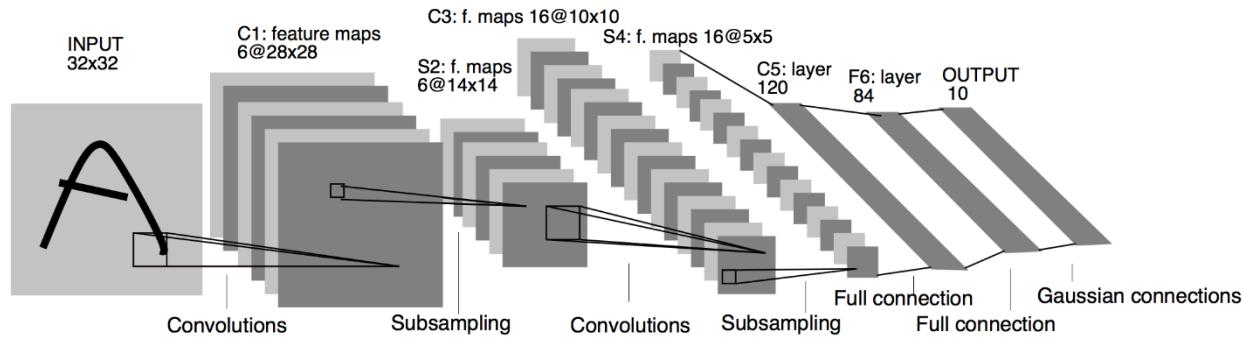
3 Further Neural Networks

03. Further Neural Networks

Deep Learning Architectures



Example: LeNet-5 (1998)



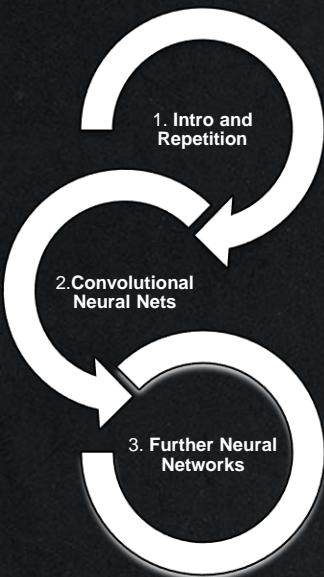
Key features:

- Convolutional for spatial features
- Subsampling using average pooling
- Non-linearity: tanh
- Sparse connectivity between S2 and C3 (efficiency, robustness)
- MLP as final classifier
- Sequence: con, pooling, non-linearity
- Foundation for many other architectures

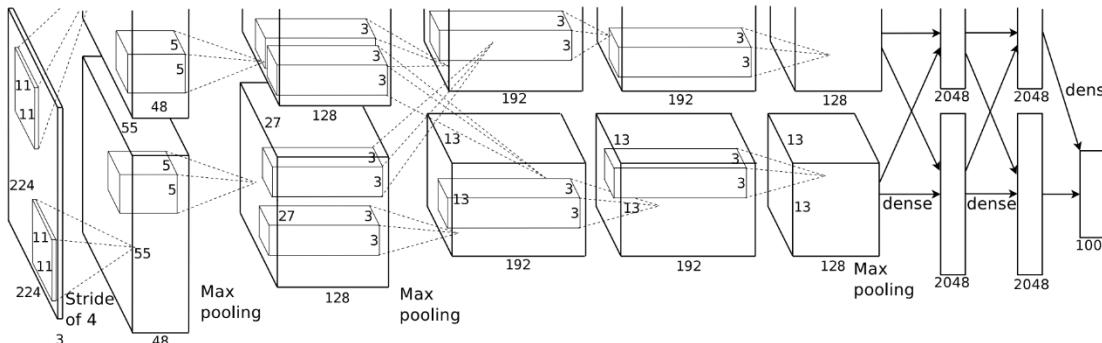
<https://medium.com/@pechyonkin/key-deep-learning-architectures-lenet-5-6fc3c59e6f4>

03. Further Neural Networks

Deep Learning Architectures



Example: AlexNet (2012)



Key features:

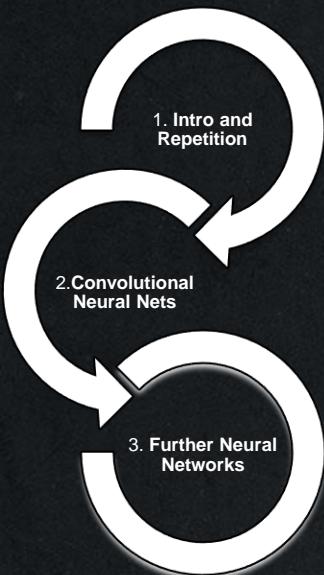
- 8 Layers
- Use of GPU(s) to reduce training time
- Overlapping max pooling
- Non-linearity: ReLu

Winner of the ImageNet 2012 challenge

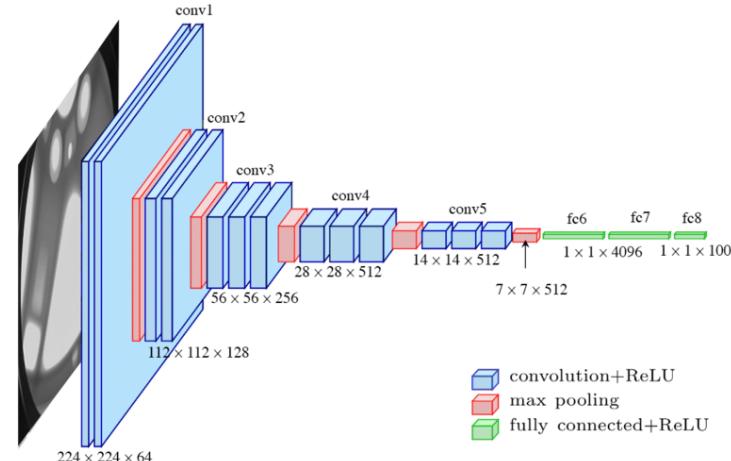
[https://paperswithcode.com/
method/alexnet#](https://paperswithcode.com/method/alexnet#)

03. Further Neural Networks

Deep Learning Architectures



Example: VGG Network (Visual Geometry Group – University of Oxford)



Key features:

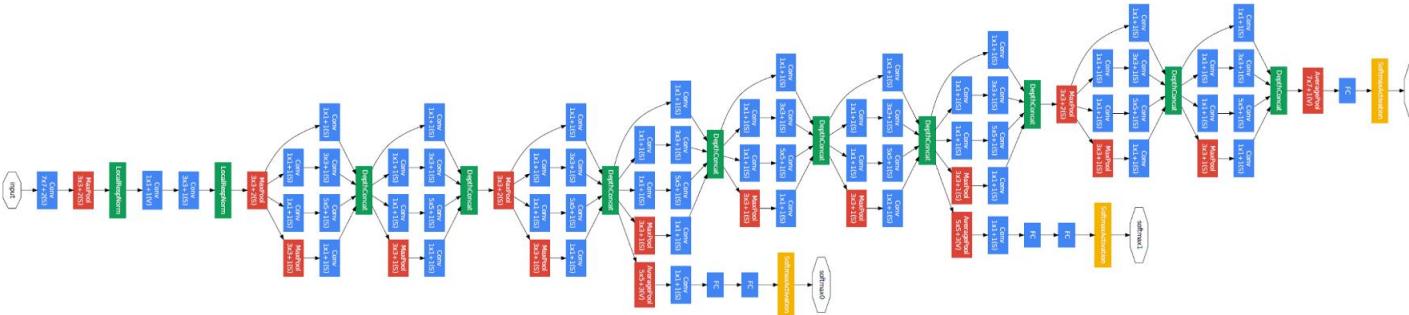
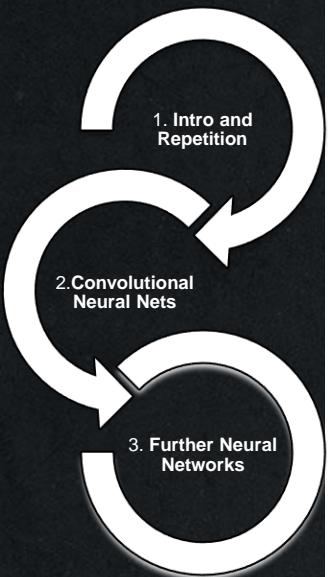
- Small kernel sizes in each convolution
- Combination of multiple smaller kernels emulate larger receptive fields
- 16/19 layers, max pooling between some layers
- Learning procedure similar to AlexNet
- Hard to train (in practice: train with shallower network)

https://www.researchgate.net/figure/Fig-A1-The-standard-VGG-16-network-architecture-as-proposed-in-32-Note-that-only_fig3_322512435

03. Further Neural Networks

Deep Learning Architectures

Example: GoogleNet (Inception-v-1)



Goal:

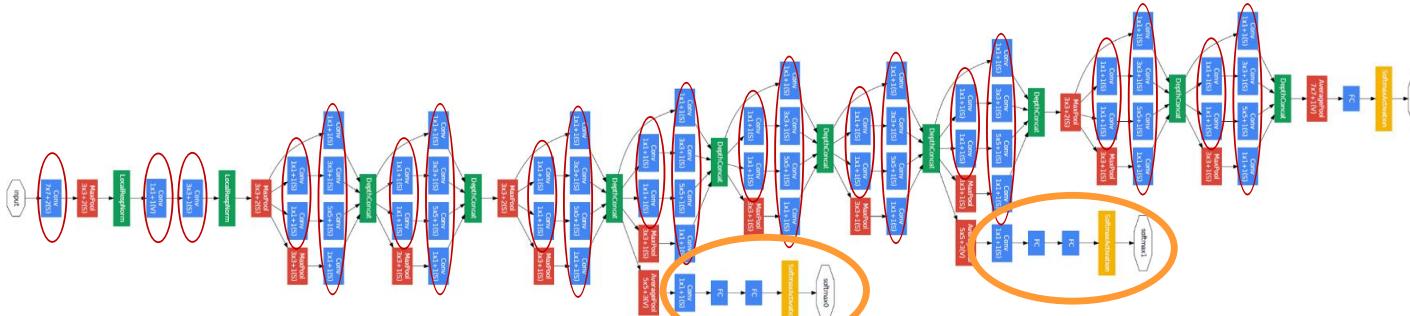
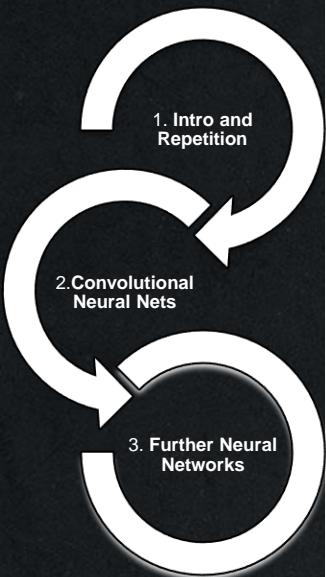
- Network design with embedded hardware in mind
=> Maximum 1.5 billion MAD (multiply-add) operations at inference time

<https://www.analyticsvidhya.com/blog/2018/10/understanding-inception-network-from-scratch/>

03. Further Neural Networks

Deep Learning Architectures

Example: GoogleNet (Inception-v-1)



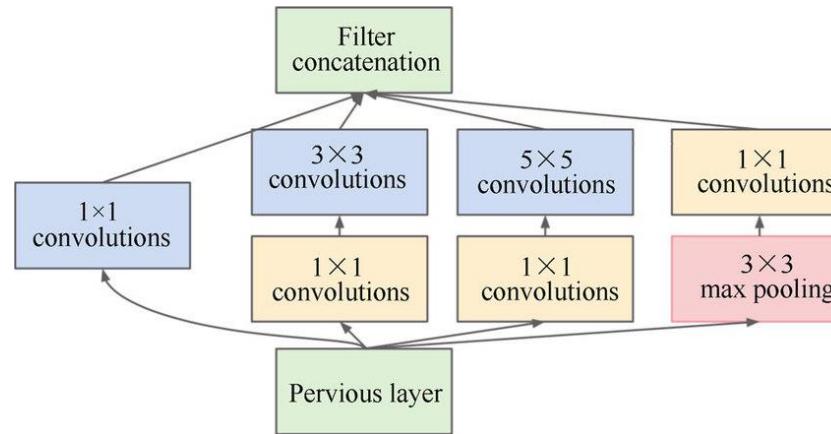
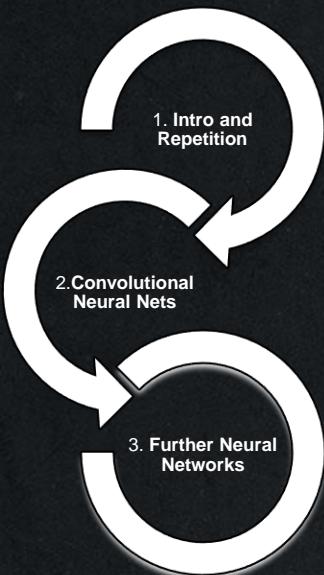
<https://www.analyticsvidhya.com/blog/2018/10/understanding-inception-network-from-scratch/>

Key features:

- 22 layers + global average pooling as final layer
- Auxiliary classifiers (only at training): error weighted by 0.3 added to global error (additional regularization)
- No fully connected layers (except for linear layer and auxiliary networks)
- Inception modules

03. Further Neural Networks

Deep Learning Architectures



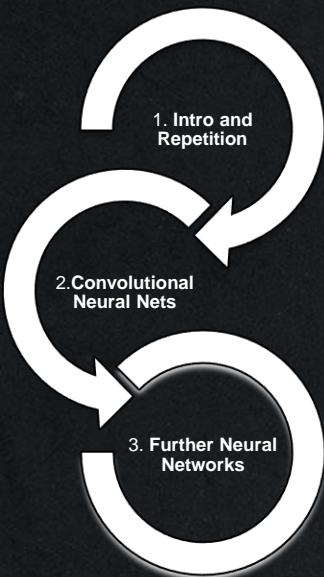
Inception Module

- derived from NiN concept
 - parallel filter combinations (split-transform-merge strategy)
 - 1×1 filter serve as „bottleneck“ layers
- ⇒ representational power of large and dense layers but with much lower computational complexity

https://www.researchgate.net/figure/inception-module-of-GooLeNet-This-figure-is-from-the-original-paper-10_fig3_312515254

03. Further Neural Networks

Deep Learning Architectures



Example: GoogleNet (Inception-v-1)

Bottleneck Layer

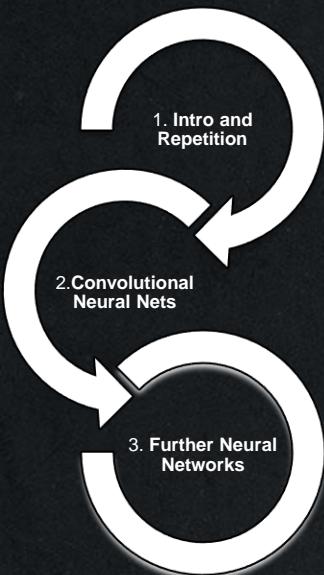
- Features are correlated, redundancy can be removed by 1×1 filters

Example:

- Before: 256 input feature maps, 256 output feature maps, 3×3 convolutions
 $\Rightarrow 256 \times 3 \times 3 \times 256 \quad 600k \text{ MAD}$
- Instead: reduce number of feature maps that have to be convolved, e.g. 64:
 - $\Rightarrow 256 \times 1 \times 1 \times 64 \quad 16k \text{ MAD}$
 - $\Rightarrow 64 \times 3 \times 3 \times 64 \quad 36k \text{ MAD}$
 - $\Rightarrow 64 \times 1 \times 1 \times 256 \quad 16k \text{ MAD}$
 - $\Rightarrow 70k \text{ MAD}$

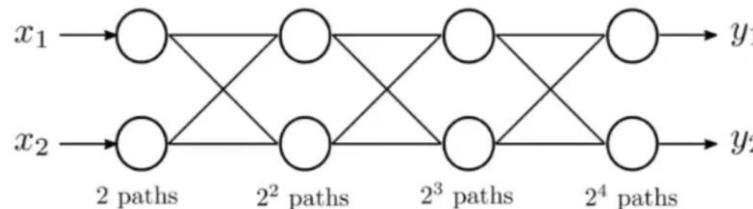
03. Further Neural Networks

Deep Learning Architectures

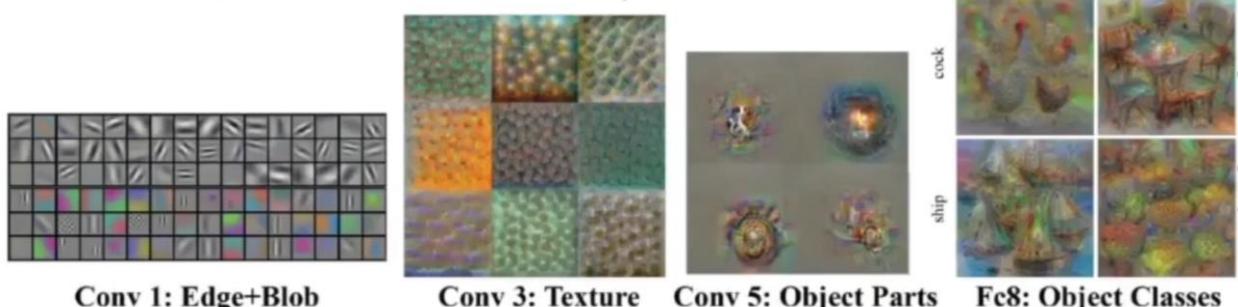


Advantages of deeper Networks

- Exponential feature reuse



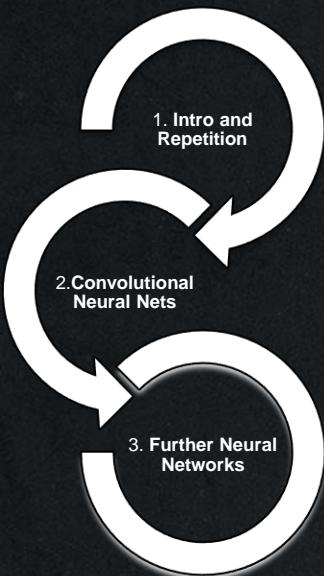
- Increasingly abstract features



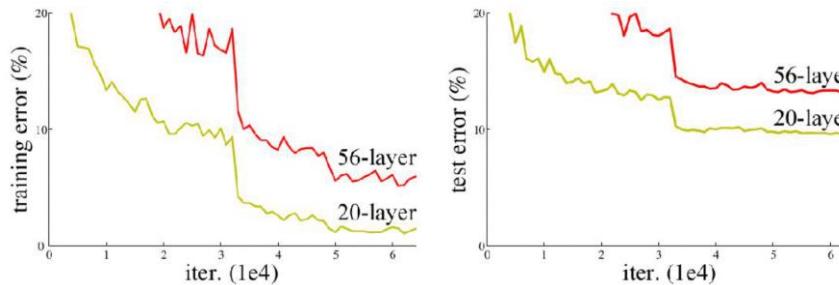
Source: http://vision03.csail.mit.edu/cnn_art/index.html

03. Further Neural Networks

Deep Learning Architectures



Advantages of deeper Networks



Deeper models tend to have **higher training and validation error** than shallower model

⇒ Not just caused by overfitting

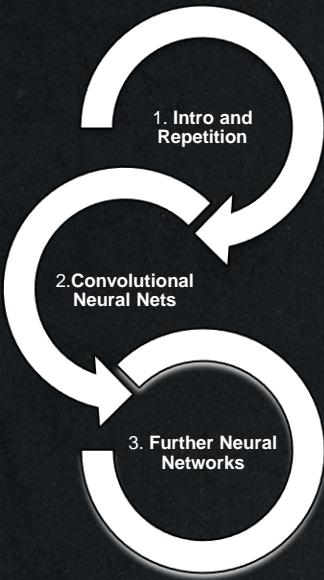
Possible reasons:

- Vanishing gradient problem
- ReLu (or successor)
- Internal co-variate shifts
 - Batch normalisation
 - ELU / SELU

Degradation problem: poor propagation of activations and gradients

03. Further Neural Networks

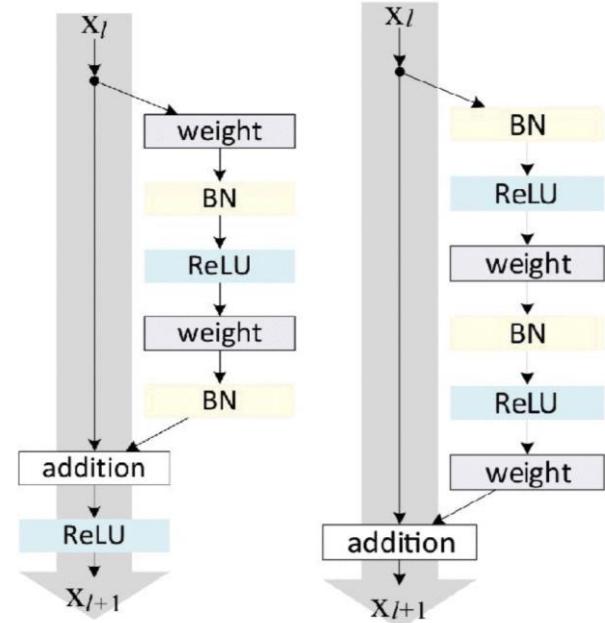
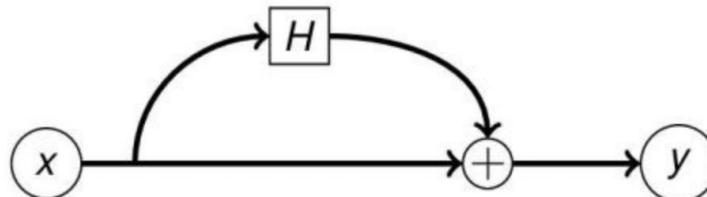
Deep Learning Architectures



Residual Units: simplify “identity solution”

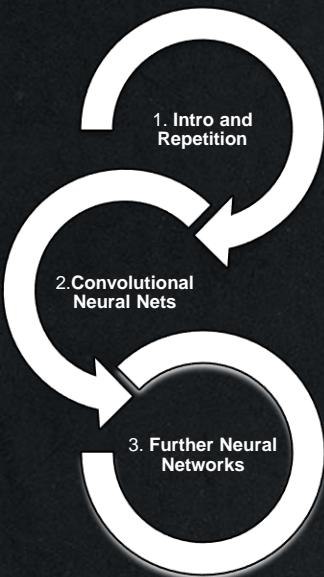
- Non-residual nets: learn mapping $F(x)$
- instead: learn residual mapping:

$$H(x) = F(x) - x \Leftrightarrow F(x) = H(x) + x$$



03. Further Neural Networks

Deep Learning Architectures

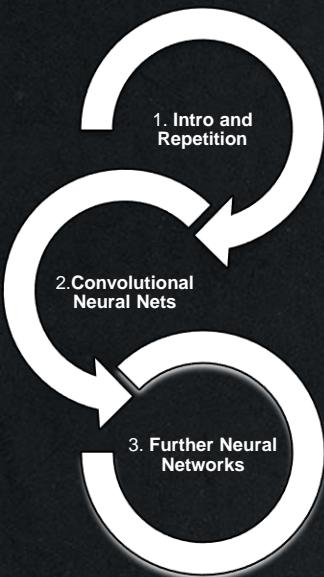


Recurrent Neural Networks

- So far: **one** input, e.g. single image
- **Feedforward** neural network: input => processing => result
- But: lots of sequential or time-dependent signals, e.g.
 - speech / music: translation; classification
 - video: object detection; face recognition
 - sensor data: speed; temperature; energy consumption
- „Snapshots“ often not informative (single word => translation ???)
- **Temporal context** is important!!

03. Further Neural Networks

Deep Learning Architectures

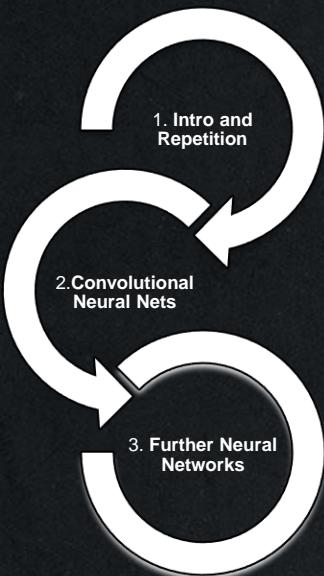


Recurrent Neural Networks

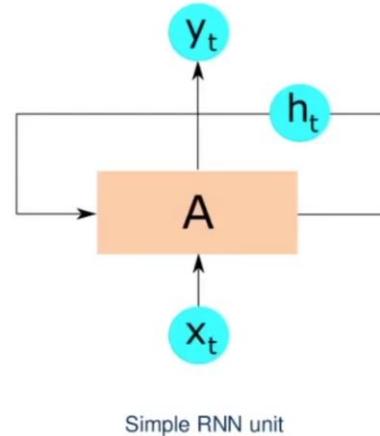
- Question: how to integrate this context into a network?
- Simple approach: feed the whole sequence to a big network: **bad idea**
 - inefficient memory usage
 - difficult / impossible to train
 - Difference between spatial and temporal dimensions?
- Better approach: model sequential behaviour within the architecture:
=> Recurrent Neural Networks (RNN)

03. Further Neural Networks

Deep Learning Architectures



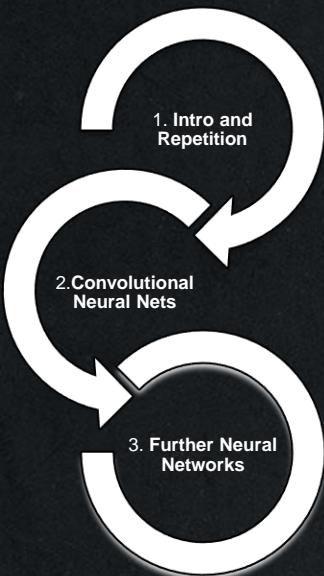
Recurrent Neural Networks



- First model in 1970s and 1980s (Hopfield Network)
- Simple **Recurrent Neural Networks (RNN)** or **Elman networks** introduced in *Finding Structure in Time* by Jeff Elman

03. Further Neural Networks

Deep Learning Architectures

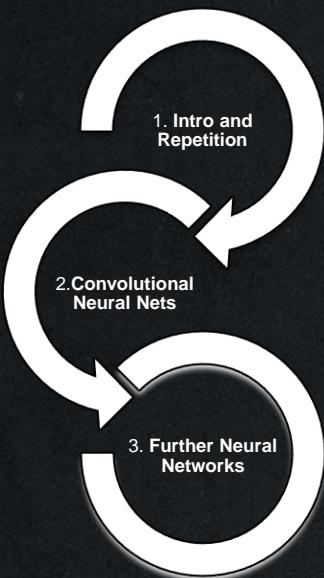


Recurrent Neural Networks – Differences to Feedforward Neural Networks

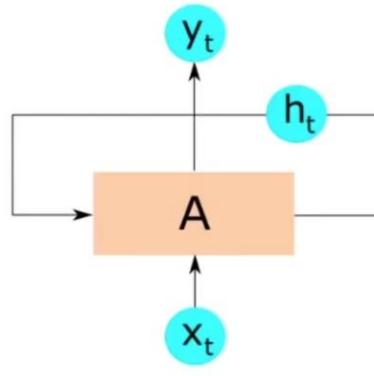
- Feedforward networks **only feed** information **forward**
- With recurrent neural networks, we can:
 - Model loops
 - Model memory and experience
 - Learn sequential relations
 - Provide continuous predictions as data comes in = real time

03. Further Neural Networks

Deep Learning Architectures



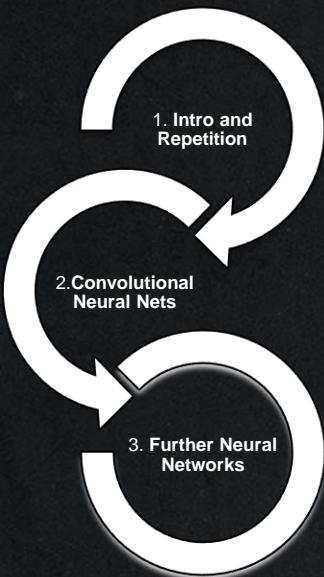
Recurrent Neural Networks – Basic Structure



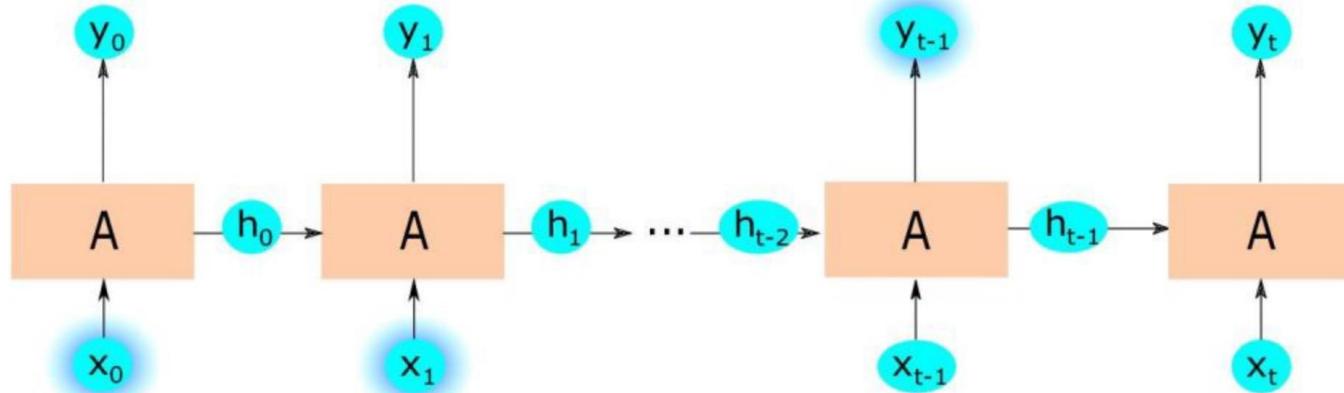
- Current input multiplied by weight
 - Additional input: hidden state of the unit
- => feedback loop: use information from present and recent past to compute output

03. Further Neural Networks

Deep Learning Architectures



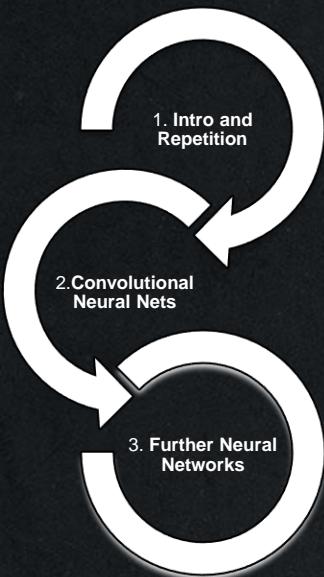
Recurrent Neural Networks – Basic Structure



- „Unfolded RNN“ unit: sequence of copies of the **same** unit (same weights)
- each unit passes hidden state as additional input to successor
- ⇒ Previous input can influence current output

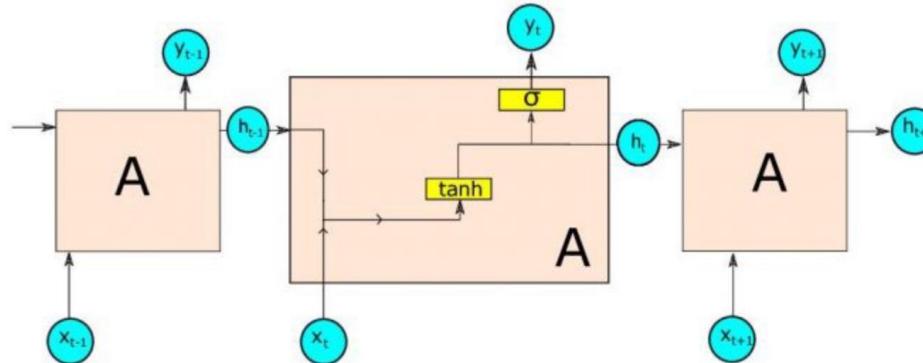
03. Further Neural Networks

Deep Learning Architectures



Recurrent Neural Networks – close-up of a basic RNN unit

- Question 1: how do we update the hidden state?
- Question 2: how do we combine input and hidden state to compute output?

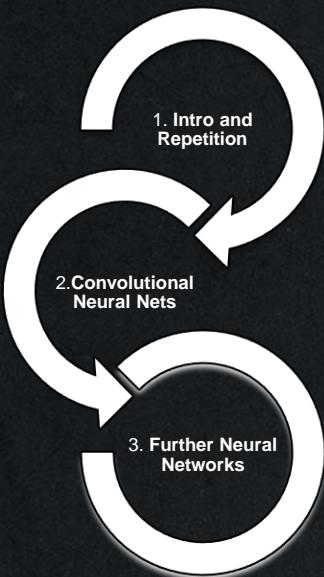


Two activation functions:

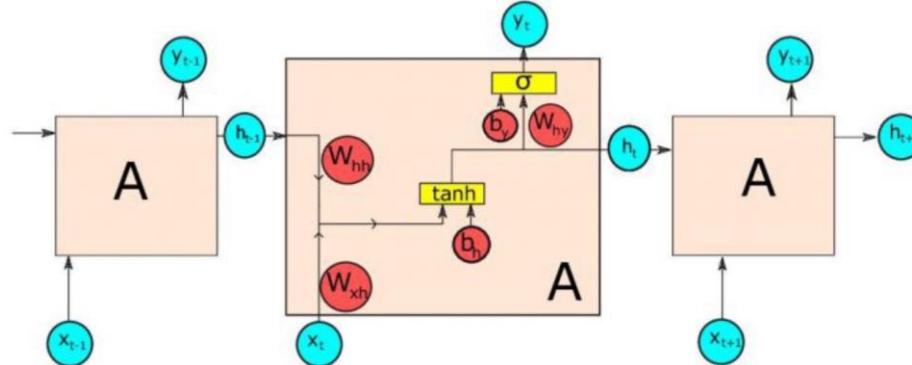
- \tanh : combination of previous state and current input
- σ : additional non-linearity for output

03. Further Neural Networks

Deep Learning Architectures



Recurrent Neural Networks – close-up: how to update the hidden state



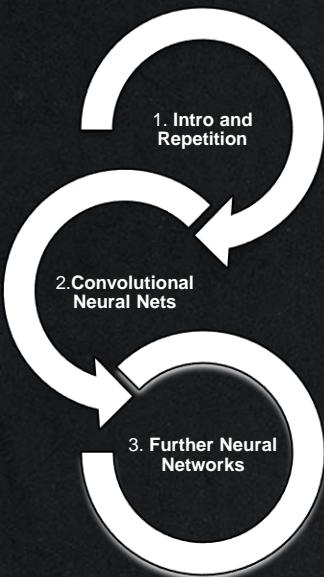
Update hidden state:

$$h_t = \tanh(W_{hh} \cdot h_{t-1} + W_{xh} \cdot x_t + b_h)$$

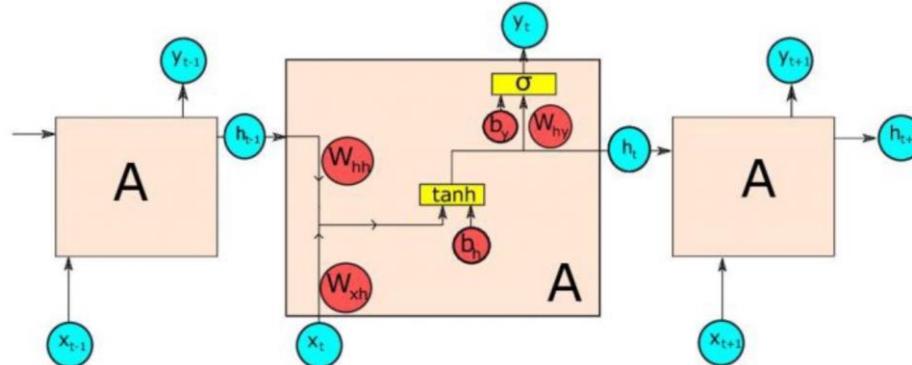
- W_{hh} : weight matrix for previous hidden state
- W_{xh} : weight matrix for current input
- b_h : update bias

03. Further Neural Networks

Deep Learning Architectures



Recurrent Neural Networks – close-up: how to update the hidden state



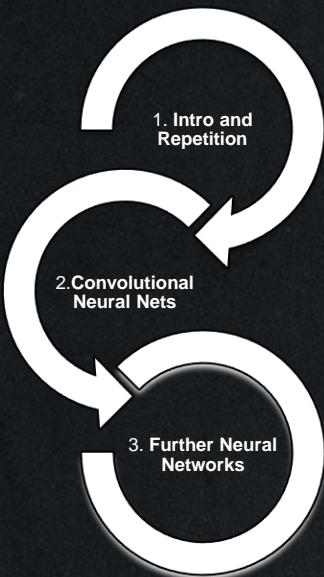
Update hidden state:

$$y_t = \sigma(W_{hy} \cdot h_t + b_y)$$

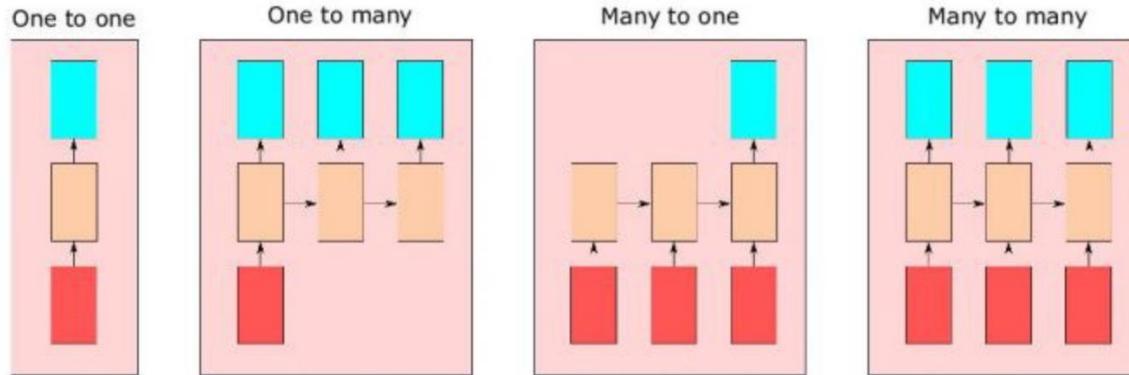
- W_{hy} : weight matrix for current hidden state
- b_h : output bias

03. Further Neural Networks

Deep Learning Architectures



Recurrent Neural Networks – Architectures

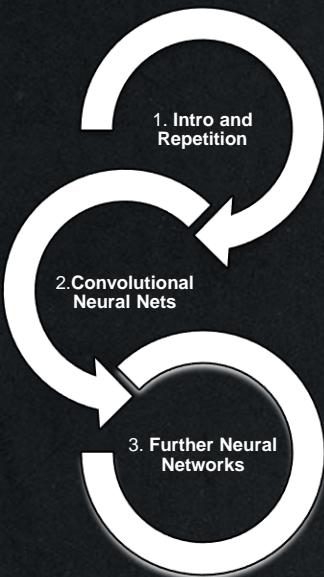


Examples:

- **one-to-one:** image classification (classic feedforward)
- **one-to-many:** image captioning
- **many-to-one:** sentiment analysis
- **many-to-many:** video classification

03. Further Neural Networks

Deep Learning Architectures



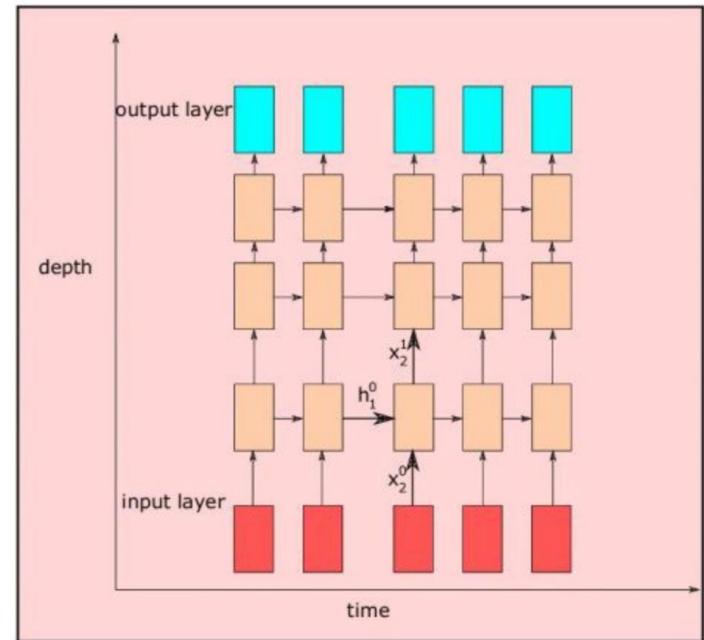
Recurrent Neural Networks – Architectures

Similar to CNN: stack multiple units for **deep RNNs**

$$h_t' = \tanh(W_{xh}^L \cdot x_t^L + W_{hh}^L \cdot h_{t-1}^L + b^L)$$

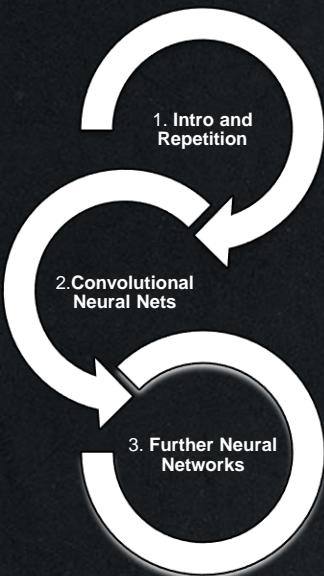
t : time point

L : layer index



03. Further Neural Networks

Deep Learning Architectures



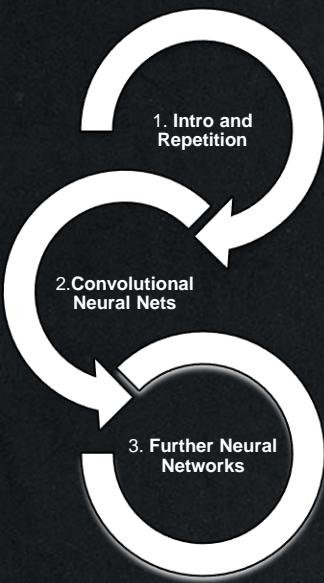
Recurrent Neural Networks – Training

Example: learn character probability distribution from input text

- Vocabulary: $\{h, e, l, o\}$
- Characters encoded as one-hot vectors, e.g. $h = (1, 0, 0, 0)$
- Train RNN on the sequence „hello“:
 - Given „h“ as first input, the network should generate sequence „hello“
 - Network needs to know previous inputs when presented with „l“:
 - do we need another „l“ or an „o“?

03. Further Neural Networks

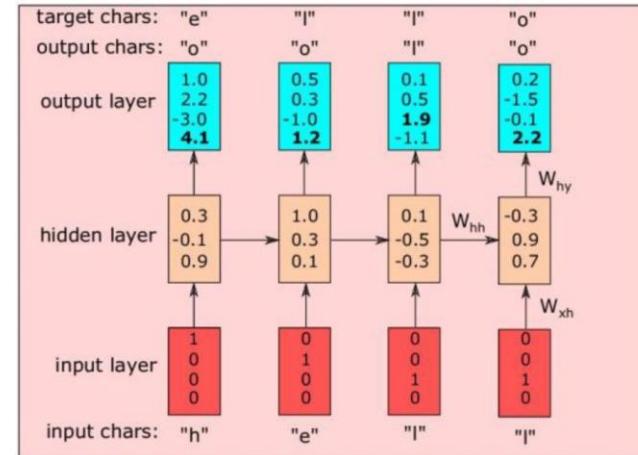
Deep Learning Architectures



Recurrent Neural Networks – Training

Example: learn character probability distribution from input text

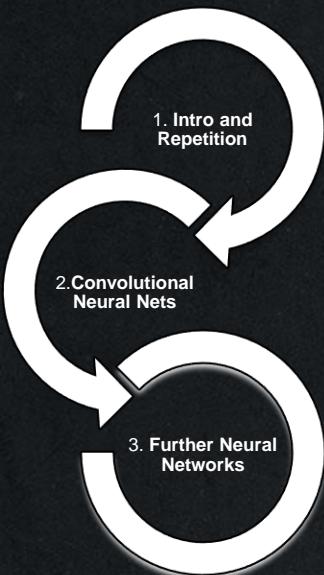
Prediction with random initialization



- Goal: maximize prediction for correct component
 - How can we now train this network?
- ⇒ Backpropagation through time (BPTT): train „unfolded“ network

03. Further Neural Networks

Deep Learning Architectures



Long Short-Term Memory Units (LSTM)

- Introduced by Hochreiter & Schmidhuber in 1997
- Designed to solve vanishing gradient and learning long-term dependencies

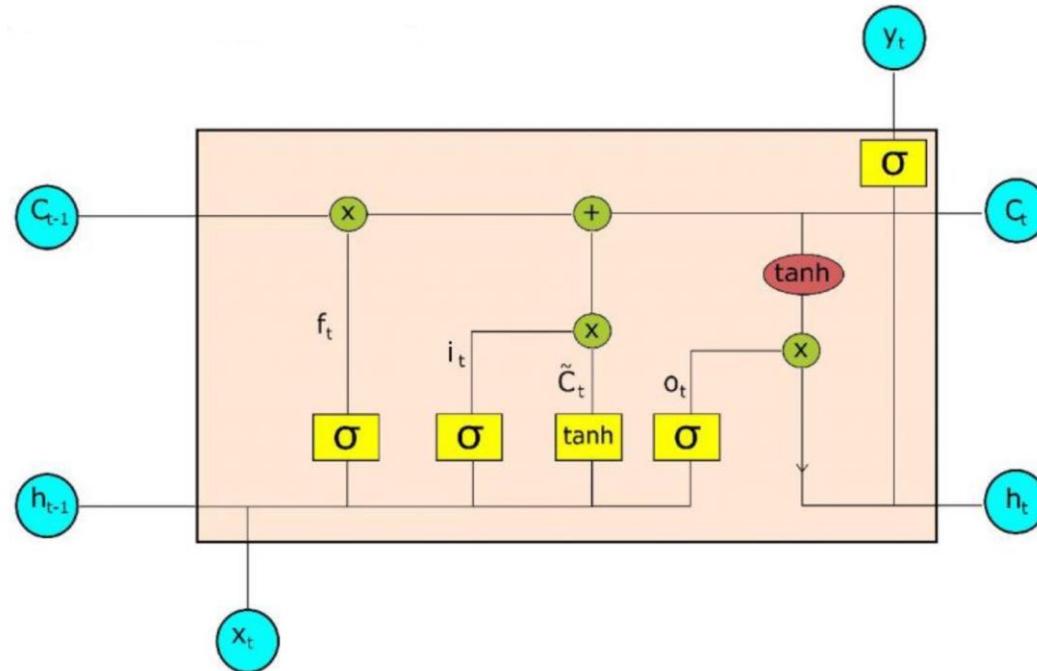
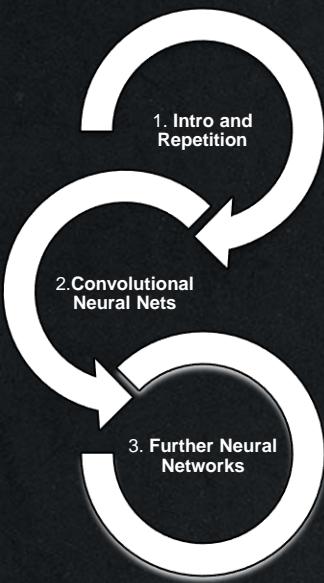
Main idea:

⇒ Introduce gates that control writing and accessing „memory“ in additional cell state

03. Further Neural Networks

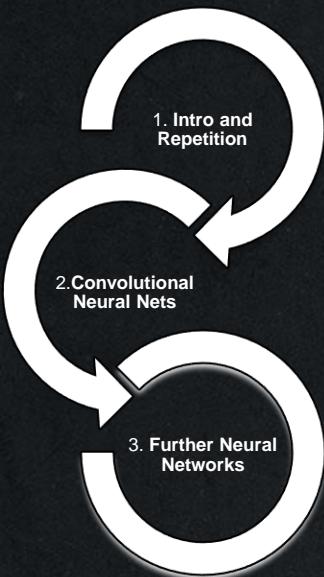
Deep Learning Architectures

Long Short-Term Memory Units (LSTM)



03. Further Neural Networks

Deep Learning Architectures



Long Short-Term Memory Units (LSTM) - Workflow

Elements of LSTM units:

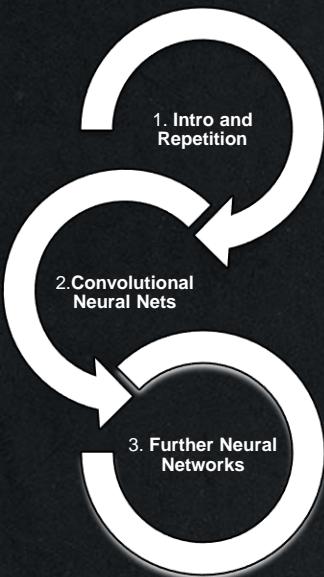
- Input x_t
- Hidden state $h_{t-1}; h_t$
- Cell state $C_{t-1}; C_t$
- Output y_t

Update of internal states in multiple steps:

- Forget Gate: forgetting old information in cell state
- Input Gate: deciding on new input for cell state
- Computing the updated cell state
- Computing the updated hidden state

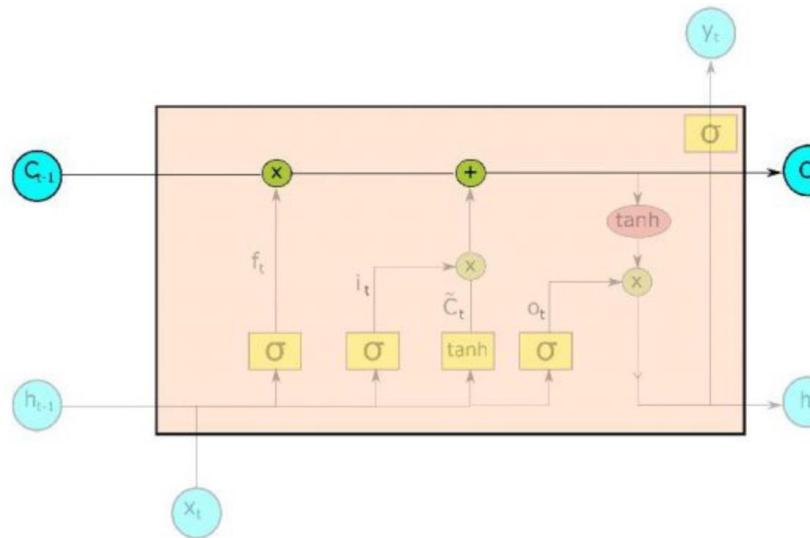
03. Further Neural Networks

Deep Learning Architectures



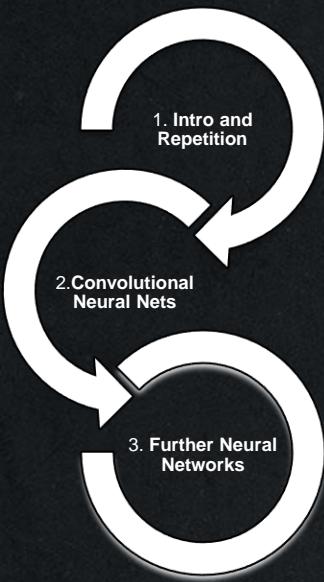
Long Short-Term Memory Units (LSTM) – Cell State

- C_t : cell state after time point t
- undergoes only linear changes: no activation function!
- C_t can flow through a unit unchanged => cell state can be constant for multiple time steps



03. Further Neural Networks

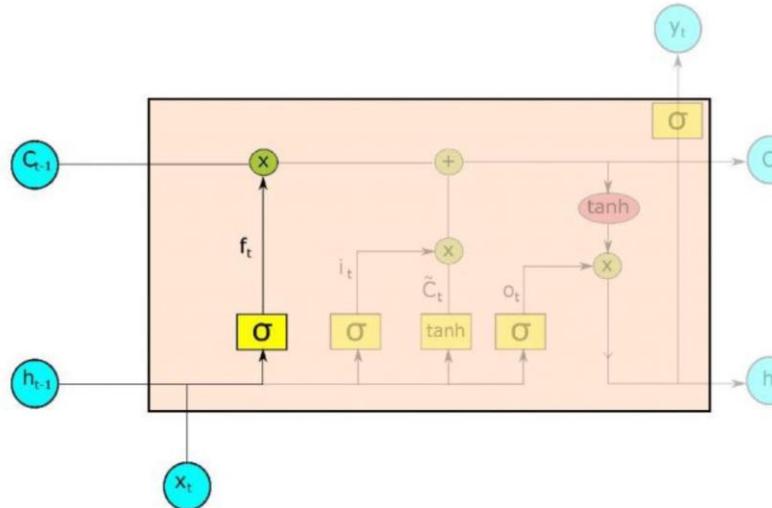
Deep Learning Architectures



Long Short-Term Memory Units (LSTM) – Forgetting old information

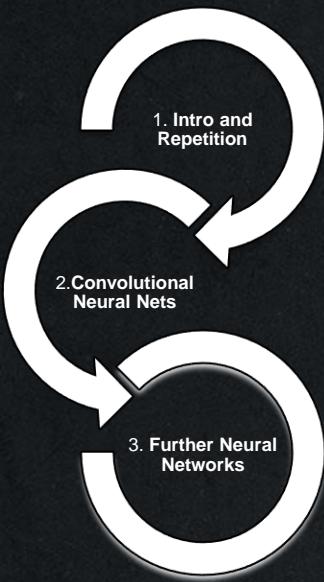
- **Key idea:** „forgetting“ and „memorizing“ information in separate steps
- controls how much of the previous cell state is forgotten

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$



03. Further Neural Networks

Deep Learning Architectures

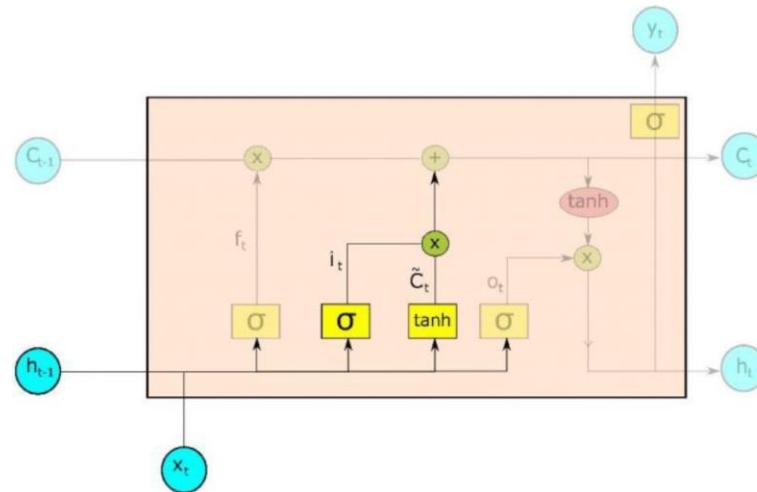


Long Short-Term Memory Units (LSTM) – Input Gate

Combination of input and hidden state on two paths:

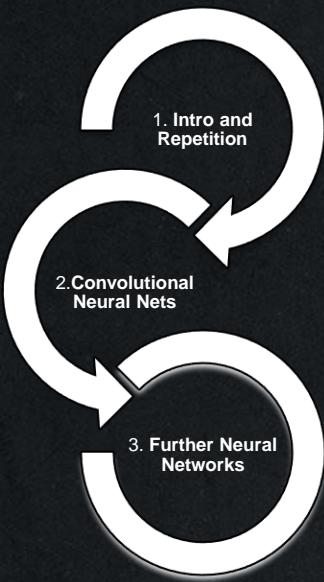
$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$



03. Further Neural Networks

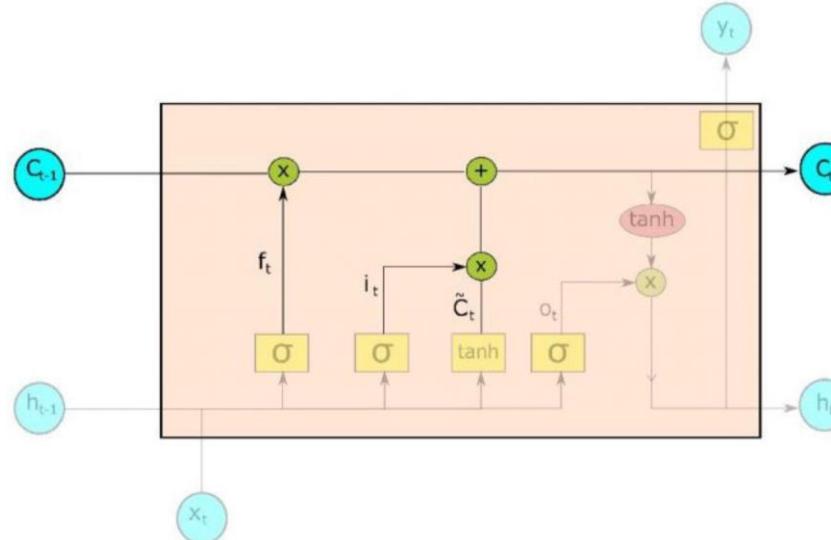
Deep Learning Architectures



Long Short-Term Memory Units (LSTM) – Updating Cell State

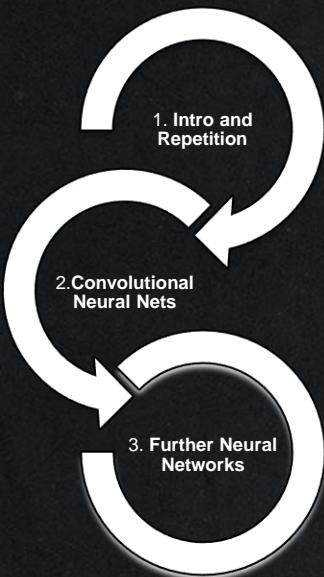
New cell state: sum of „remaining information“ from and new information from input and hidden state

$$C_t = f_t \odot C_{t-1} + i_t \odot \tilde{C}_t$$



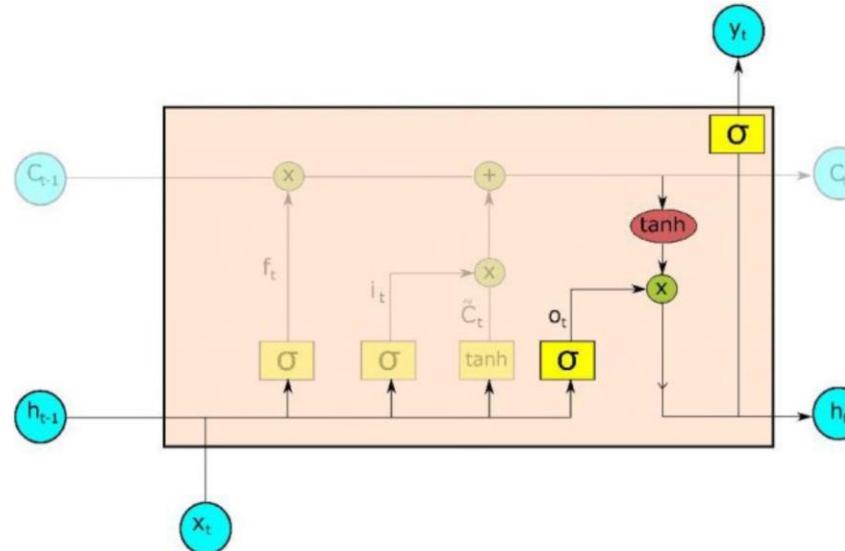
03. Further Neural Networks

Deep Learning Architectures



Long Short-Term Memory Units (LSTM) – Updating Hidden State and Output

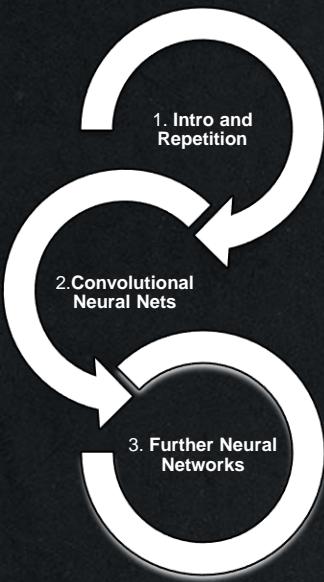
- Important: cell state and hidden state are updated **separately**
- Output y_t directly depends on the hidden state h_t



03. Further Neural Networks

Deep Learning Architectures

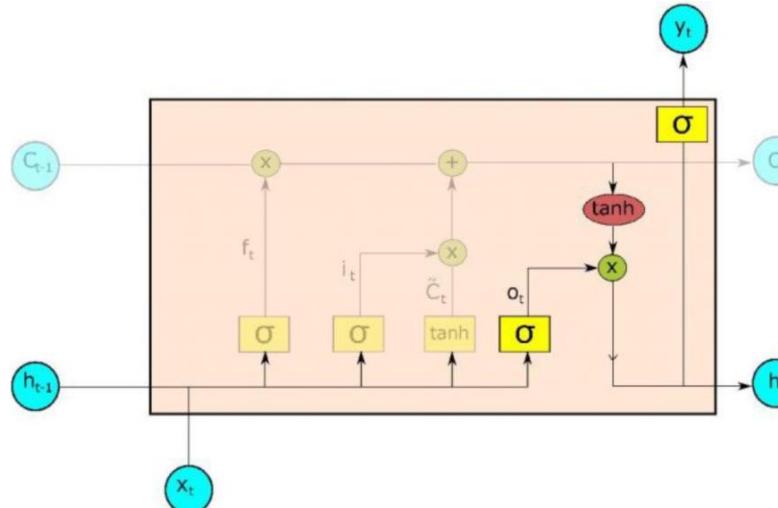
Long Short-Term Memory Units (LSTM) – Updating Hidden State and Output (cont.)



$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o)$$

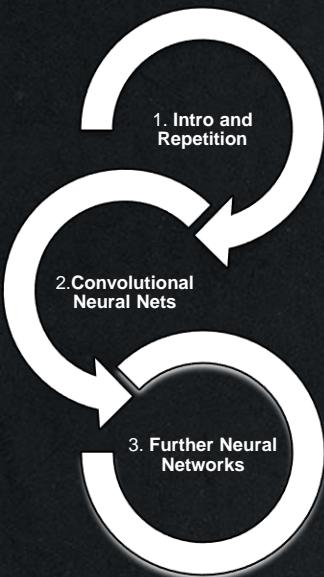
$$h_t = o_t \odot \tanh(C_t)$$

$$y_t = \sigma(h_L)$$



03. Further Neural Networks

Deep Learning Architectures



Recurrent / LSTM Neural Networks – Summary

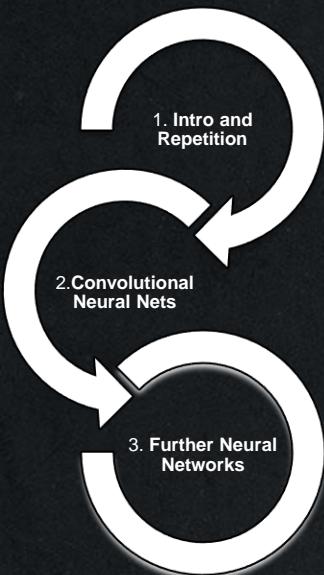
- RNN are able to directly model sequential algorithms
- Training through (truncated) backpropagation through time
- Simple units suffer extremely from exploding / vanishing gradients
- LSTM nets as improved RNN units explicitly modelling „forgetting“ / „remembering“

We haven't talked about:

- Gated Recurrent Unit (GRU)
- Neural turing machines
- Only grazed: attention + recurrent neural networks
-

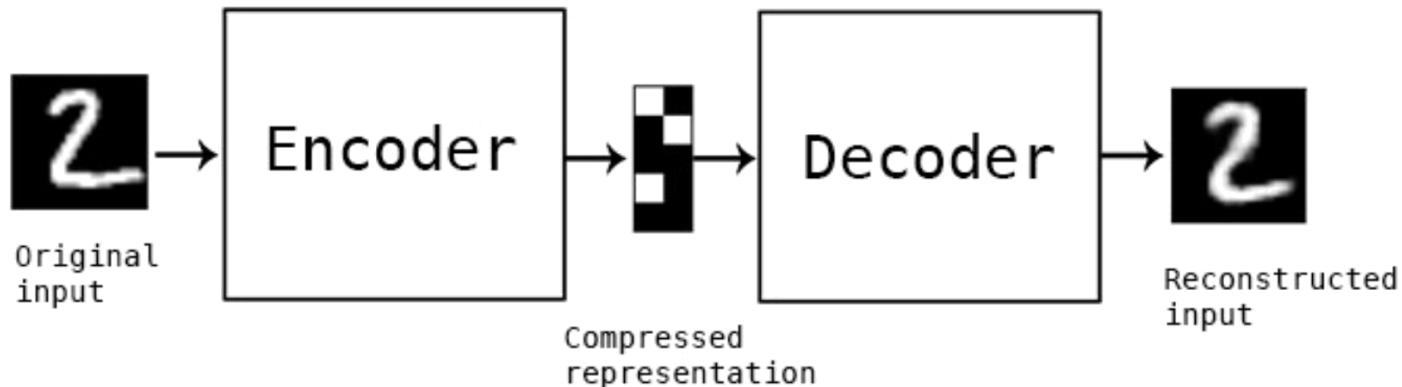
03. Further Neural Networks

Deep Learning Architectures



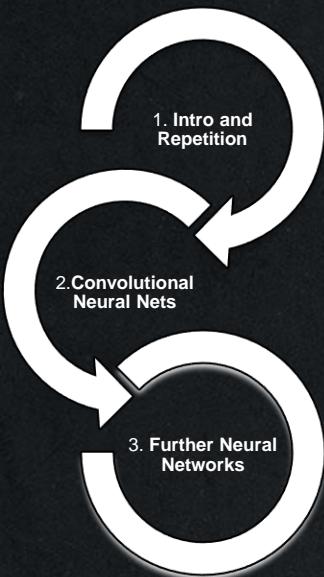
Autoencoder - Definition

- Autoencoder is an unsupervised artificial neural network that learns how to efficiently compress and encode data then learns how to reconstruct the data back from the reduced encoded representation to a representation that is as close to the original input as possible
- Autoencoder by design reduces data dimensions by learning how to ignore the noise in the



03. Further Neural Networks

Deep Learning Architectures

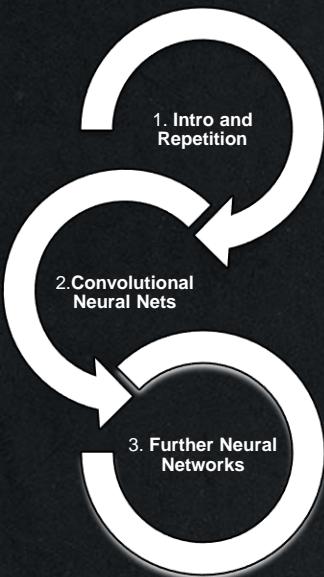


Autoencoder - Components

- 1) Encoder: In which the model learns how to reduce the input dimensions and compress the input data into an encoded representation
- 2) Bottleneck: which is the layer that contains the compressed representation of the input data. This is the lowest possible dimensions of the input data
- 3) Decoder: In which the model learns how to reconstruct the data from the encoded representation to be as close to the original input as possible
- 4) Reconstruction Loss: This is the method that measures measure how well the decoder is performing and how close the output is to the original input

03. Further Neural Networks

Deep Learning Architectures



Autoencoder - Architectures for Anomaly Detection

- There are many ways and techniques to detect anomalies and outliers
- (see)

5 Ways to Detect Outliers That Every Data Scientist Should Know (Python Code)

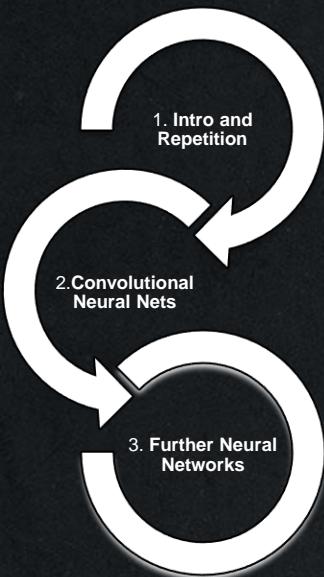
Detecting Anomalies is critical to any business either by identifying faults or being proactive. This article discusses...

[towardsdatascience.com](https://towardsdatascience.com/5-ways-to-detect-outliers-that-every-data-scientist-should-know-python-code-3e5c6f017726)



03. Further Neural Networks

Deep Learning Architectures



Autoencoder - Architectures for Anomaly Detection: Example

- If you have correlated input data, the autoencoder method will work very well because the encoding operation relies on the correlated features to compress the data
- Let's say that we have trained an autoencoder on the MNIST dataset
- Using a simple FeedForward neural network, we can achieve this by building a simple 6 layers network as below:

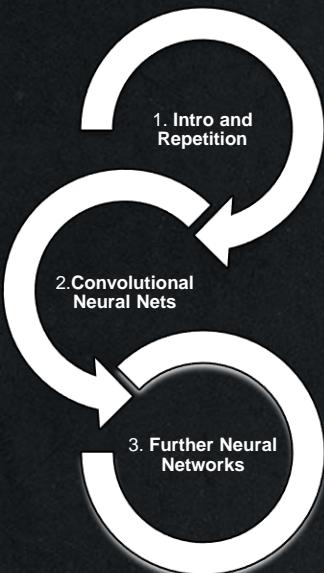
```
import numpy as np
import keras
from keras.datasets import mnist
from keras.models import Sequential, Model
from keras.layers import Dense, Input
from keras import optimizers
from keras.optimizers import Adam

(x_train, y_train), (x_test, y_test) = mnist.load_data()
train_x = x_train.reshape(60000, 784) / 255
val_x = x_test.reshape(10000, 784) / 255
```

03. Further Neural Networks

Deep Learning Architectures

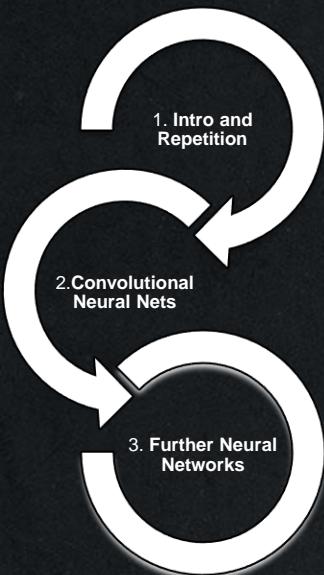
Autoencoder - Architectures for Anomaly Detection: Example



```
autoencoder = Sequential()
autoencoder.add(Dense(512, activation='elu', input_shape=(784,)))
autoencoder.add(Dense(128, activation='elu'))
autoencoder.add(Dense(10, activation='linear', name="bottleneck"))
autoencoder.add(Dense(128, activation='elu'))
autoencoder.add(Dense(512, activation='elu'))
autoencoder.add(Dense(784, activation='sigmoid'))
autoencoder.compile(loss='mean_squared_error', optimizer = Adam())
trained_model = autoencoder.fit(train_x, train_x, batch_size=1024, epochs=10, verbose=1,
validation_data=(val_x, val_x))
encoder = Model(autoencoder.input, autoencoder.get_layer('bottleneck').output)
encoded_data = encoder.predict(train_x) # bottleneck representation
decoded_output = autoencoder.predict(train_x) # reconstruction
encoding_dim = 10
```

03. Further Neural Networks

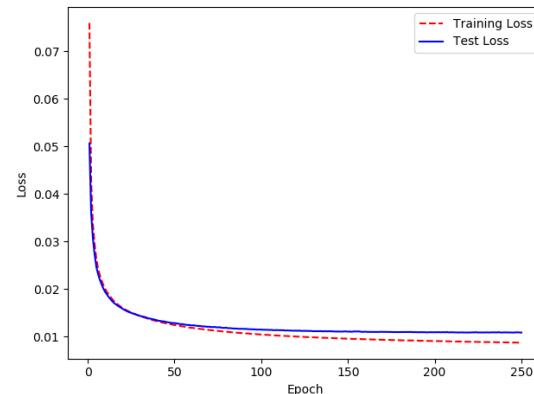
Deep Learning Architectures



Autoencoder - Architectures for Anomaly Detection: Example

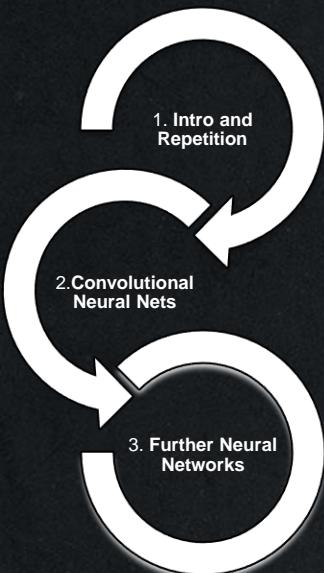
- As you can see in the output, the last reconstruction loss/error for the validation set is 0.0193
- Now, if I pass any normal image from the MNIST dataset, the reconstruction loss will be very low (< 0.02)
- BUT if I tried to pass any other different image (outlier or anomaly), we will get a high reconstruction loss value because the network failed to reconstruct the image/input that is considered an anomaly

```
# return the decoder
encoded_input = Input(shape=(encoding_dim,))
decoder = autoencoder.layers[-3](encoded_input)
decoder = autoencoder.layers[-2](decoder)
decoder = autoencoder.layers[-1](decoder)
decoder = Model(encoded_input, decoder)
```



03. Further Neural Networks

Deep Learning Architectures



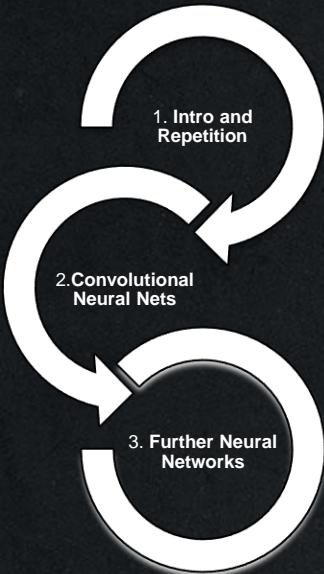
Autoencoder - Architectures for Anomaly Detection: Example

- As you can see in the output, the last reconstruction loss/error for the validation set is 0.0193
- Now, if I pass any normal image from the MNIST dataset, the reconstruction loss will be very low (< 0.02)
- BUT if I tried to pass any other different image (outlier or anomaly), we will get a high reconstruction loss value because the network failed to reconstruct the image/input that is considered an anomaly

```
# return the decoder
encoded_input = Input(shape=(encoding_dim,))
decoder = autoencoder.layers[-3](encoded_input)
decoder = autoencoder.layers[-2](decoder)
decoder = autoencoder.layers[-1](decoder)
decoder = Model(encoded_input, decoder)
```

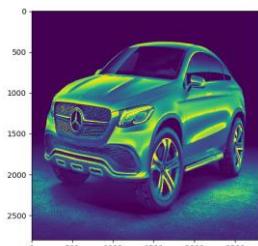
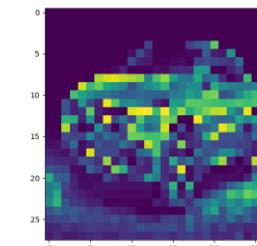
03. Further Neural Networks

Deep Learning Architectures



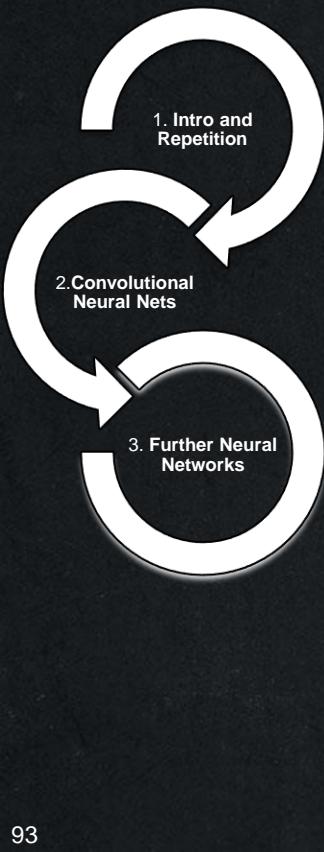
Autoencoder - Architectures for Anomaly Detection: Example

- Now, let's do some anomaly detection
- The code below uses two different images to predict the anomaly score (reconstruction error) using the autoencoder network we trained above
- The first image is from the MNIST and the result is **5.43**
- This means that the image is not an anomaly
- The second image is an image of a car, which is a completely random image that doesn't belong to the training dataset and the results were: **2362.46**
- This high error means that the image is an anomaly
- The same concept applies to any type of dataset



03. Further Neural Networks

Deep Learning Architectures



Autoencoder – Design of Bridges



sustainable concrete bridge over a river



All

Images

Videos

Shopping

News

More

Tools

drop an image of your design project location here



Drag an image here or [upload a file](#)

OR

Paste image link

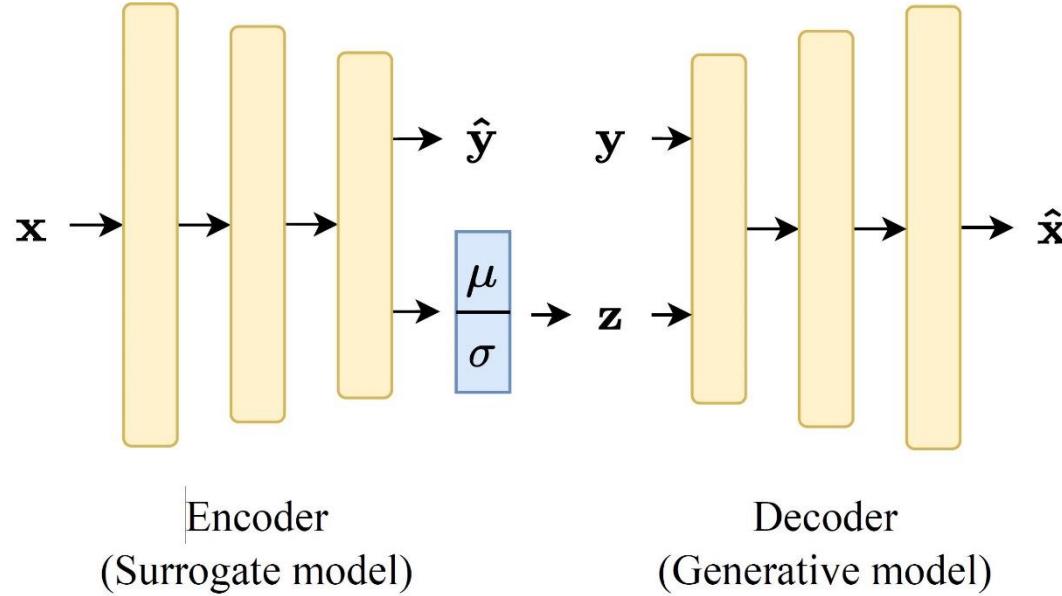
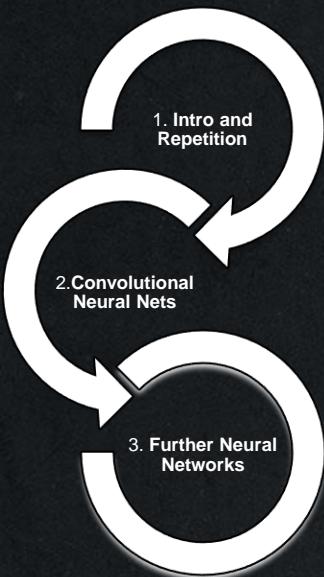
Search



03. Further Neural Networks

Deep Learning Architectures

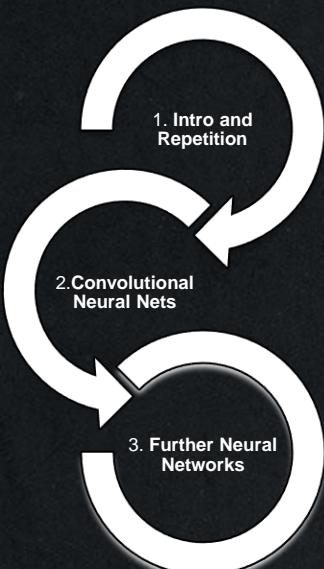
Autoencoder – Design of Bridges



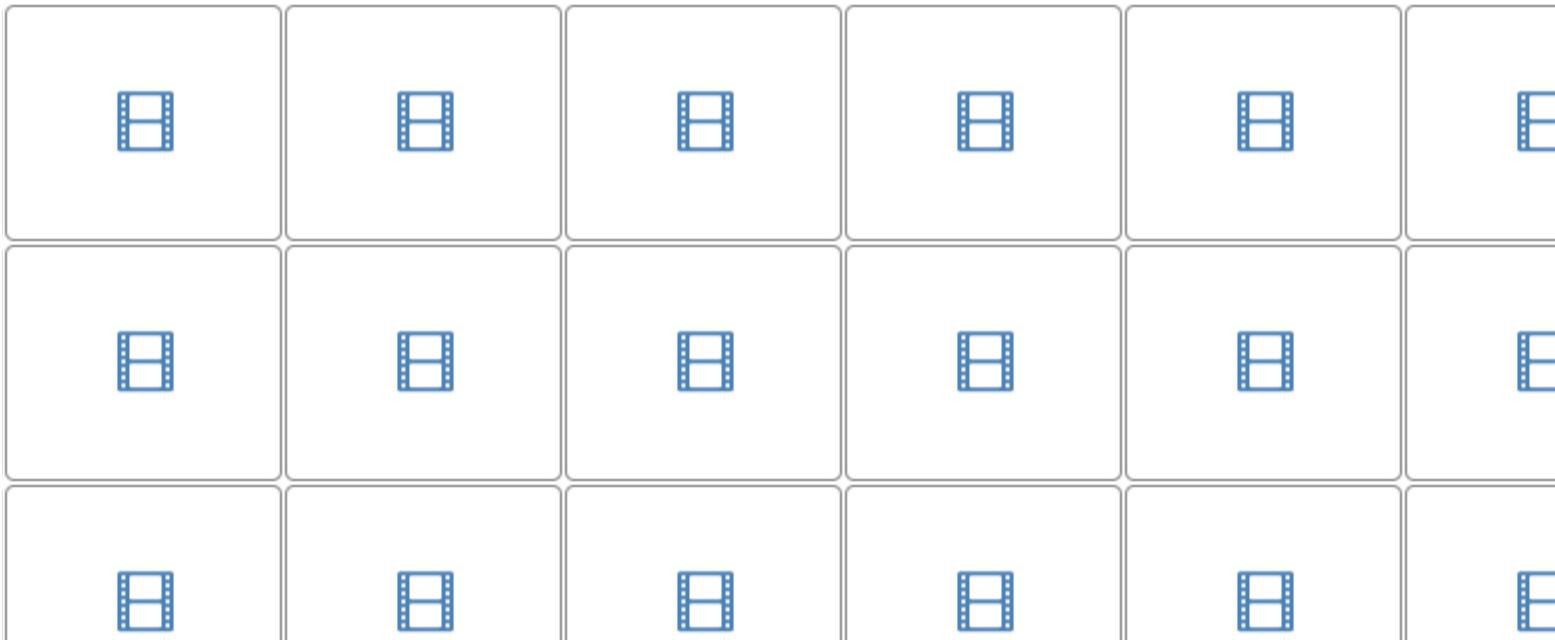
https://mkrausai.github.io/research/01_SciML/01_BH_PedestrianBridge_XAI/

03. Further Neural Networks

Deep Learning Architectures



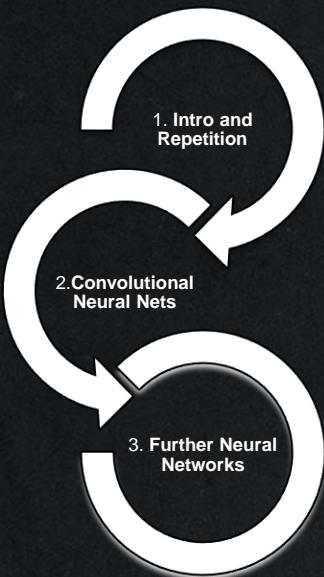
Autoencoder – Design of Bridges



https://mkrausai.github.io/research/01_SciML/01_BH_PedestrianBridge_XAI/

03. Further Neural Networks

Deep Learning Architectures



Generative Adversarial Networks (GANs)

Let's play a game:

Discriminator



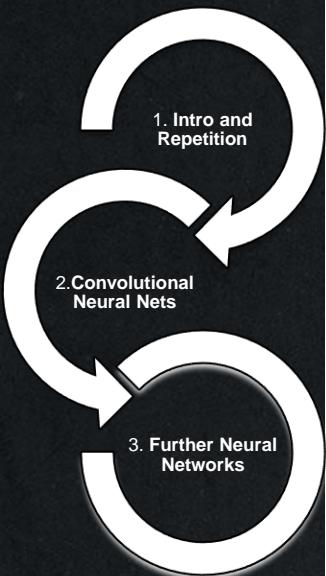
Generator



03. Further Neural Networks

Deep Learning Architectures

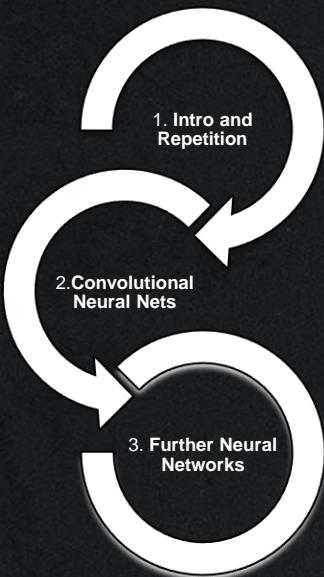
Generative Adversarial Networks (GANs)



<https://www.youtube.com/watch?v=djsEKYuiRFE&feature=youtu.be>

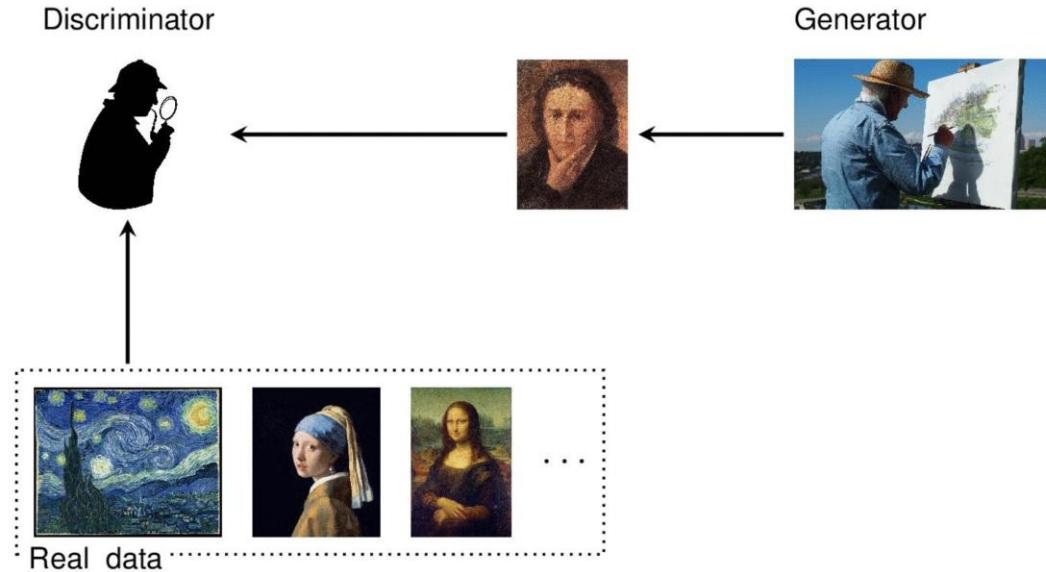
03. Further Neural Networks

Deep Learning Architectures



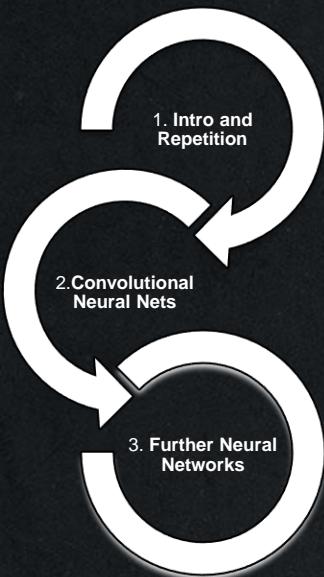
Generative Adversarial Networks (GANs)

Let's play a game:



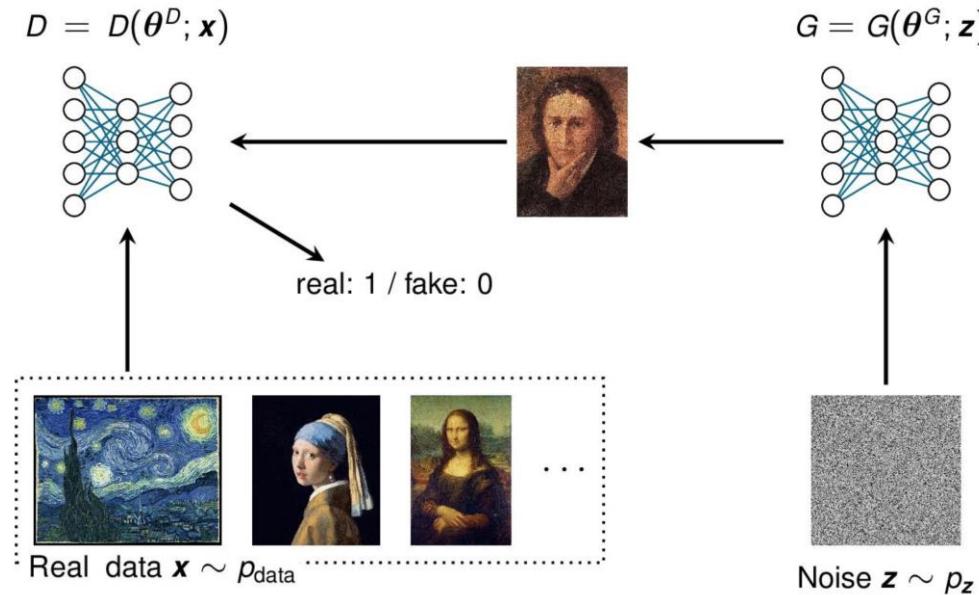
03. Further Neural Networks

Deep Learning Architectures



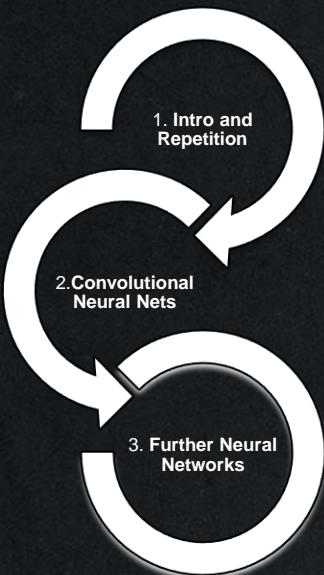
Generative Adversarial Networks (GANs)

Let's play a game:



03. Further Neural Networks

Deep Learning Architectures



Generative Adversarial Networks (GANs)

Alternate between:

- 1) Train D: minimize

$$L^D(\theta^D, \theta^G) = -\mathbb{E}_{x \sim p_{data}} \log D(x) - \mathbb{E}_{z \sim p_z} \log(1 - D(G(z)))$$

real *fake*

=> Trained to distinguish real data samples from fake ones

- 2) Train G: minimize $L^G = -L^D$

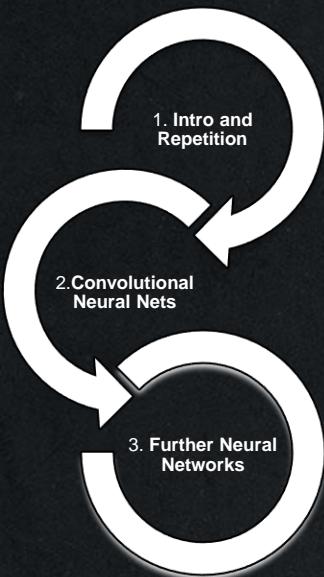
Generator minimizes log-probability of the discriminator being correct

=> Trained to generate data domain images to fool D

- Optional: run k steps of one player for every step of the other player
- Equilibrium is a saddle point of the discriminator loss

03. Further Neural Networks

Deep Learning Architectures



Generative Adversarial Networks (GANs) - Training

- assume: both densities are nonzero everywhere

- solve for:

$$\frac{\partial L^D}{\partial D(x)} = 0$$

- optimal: $D^*(x)$ for any $p_{data}(x)$ and $p_{model}(x)$:

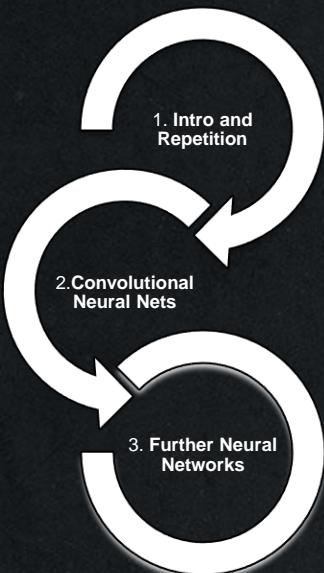
=> “just” theoretical and unachievable

$$D^*(x) = \frac{p_{data}(x)}{p_{data}(x) + p_{model}(x)}$$

- GAN’s key approximation mechanism
- GANs use supervised learning to estimate this ratio
- under- / overfitting

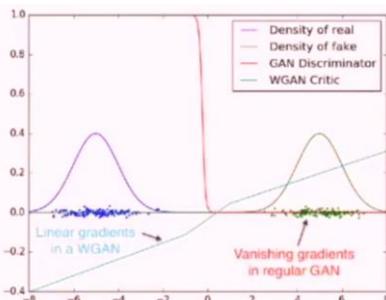
03. Further Neural Networks

Deep Learning Architectures



Generative Adversarial Networks (GANs) - Training

- Non-Saturation Game: modify Generator's Loss to maximize prob. of D being mistaken
- Feature matching loss: modify Generator's Loss to match expected feature values
- Wasserstein Loss: learn discriminator D that maximizes discrepancy between real and fake and whose gradient remains beyond a limit



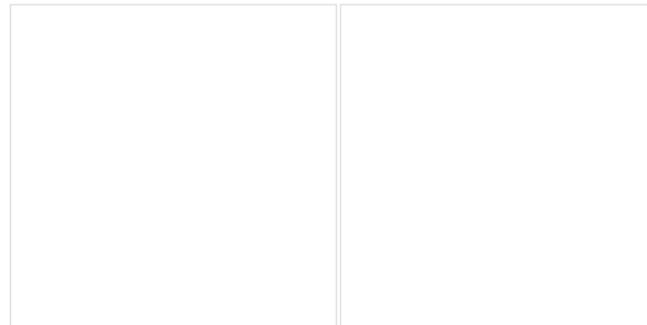
- ⇒ Bounded gradient
- ⇒ Helps countering vanishing gradients in D

03. Further Neural Networks

Deep Learning Architectures

Generative Adversarial Networks (GANs) - applied to AEC

Use GANs for floor plan generation:



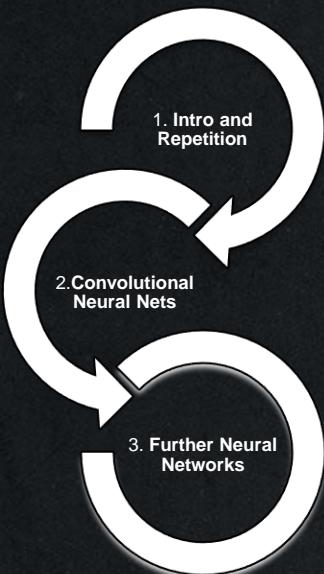
Living Room Bedroom Bathroom Corridor Kitchen Closet

Void Floor Entr. Window

Cancel Clear Transfer Random

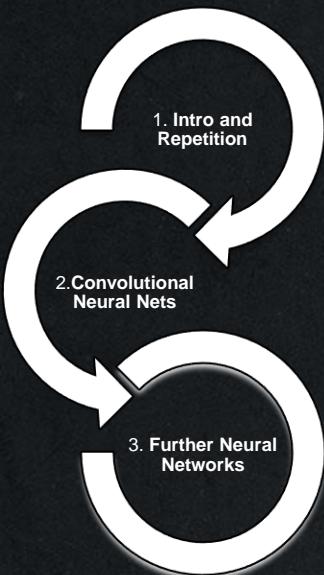
<https://developer.nvidia.com/blog/archigan-generative-stack-apartment-building-design/>

References



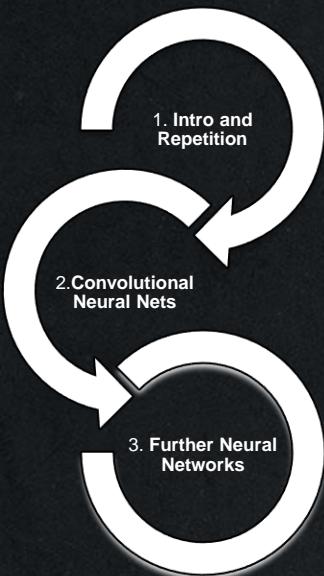
- [1] Goodfellow, I.; Bengio, Y.; Courville, A. (2016) *Deep Learning*, MIT press
- [2] Goulet, J. (2020) *Probabilistic Machine Learning for Civil Engineers*, MIT press
- [3] Bishop, C. (2006) *Pattern Recognition and Machine Learning*, Springer
- [4] Murphy, K. (2012) *Machine Learning: A Probabilistic Perspective*, MIT press
- [5] <https://stanford.edu/~shervine/teaching/cs-229/cheatsheet-machine-learning-tips-and-tricks>
- [6] David Silver, Julian Schrittwieser, Karen Simonyan, et al. "Mastering the game of go without human knowledge". In: Nature 550.7676 (2017), p. 354.
- [7] David Silver, Thomas Hubert, Julian Schrittwieser, et al. "Mastering Chess and Shogi by Self-Play with a General Reinforcement Learning Algorithm". In: arXiv preprint arXiv:1712.01815 (2017).
- [8] M. Aubreville, M. Krappmann, C. Bertram, et al. "A Guided Spatial Transformer Network for Histology Cell Differentiation". In: ArXiv e-prints (July 2017). arXiv: 1707.08525 [cs.CV].
- [9] David Bernecker, Christian Riess, Elli Angelopoulou, et al. "Continuous short-term irradiance forecasts using sky images". In: Solar Energy 110 (2014), pp. 303–315.
- [10] Alex Krizhevsky, Ilya Sutskever, and Geoffrey E Hinton. "ImageNet Classification with Deep Convolutional Neural Networks". In: Advances in Neural Information Processing Systems 25. Curran Associates, Inc., 2012, pp. 1097–1105.

References



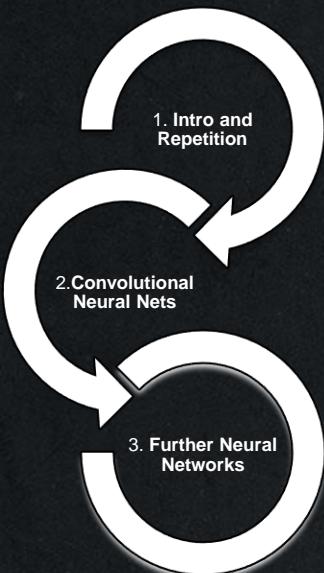
- [11] Joseph Redmon, Santosh Kumar Divvala, Ross B. Girshick, et al. "You Only Look Once: Unified, Real-Time Object Detection". In: CoRR abs/1506.02640 (2015).
- [12] J. Redmon and A. Farhadi. "YOLO9000: Better, Faster, Stronger". In: ArXiv e-prints (Dec. 2016). arXiv: 1612.08242 [cs.CV].
- [13] Joseph Redmon and Ali Farhadi. "YOLOv3: An Incremental Improvement". In: arXiv (2018).
- [14] Frank Rosenblatt. The Perceptron—a perceiving and recognizing automaton. 85–460–1. Cornell Aeronautical Laboratory, 1957.
- [15] Olga Russakovsky, Jia Deng, Hao Su, et al. "ImageNet Large Scale Visual Recognition Challenge". In: International Journal of Computer Vision 115.3 (2015), pp. 211–252.
- [16] David Silver, Aja Huang, Chris J. Maddison, et al. "Mastering the game of Go with deep neural networks and tree search". In: Nature 529.7587 (Jan. 2016), pp. 484–489.
- [17] S. E. Wei, V. Ramakrishna, T. Kanade, et al. "Convolutional Pose Machines". In: CVPR. 2016, pp. 4724–4732.
- [18] Tobias Würfl, Florin C Ghesu, Vincent Christlein, et al. "Deep learning computed tomography". In: International Conference on Medical Image Computing and Computer-Assisted Springer International Publishing. 2016, pp. 432–440.

References



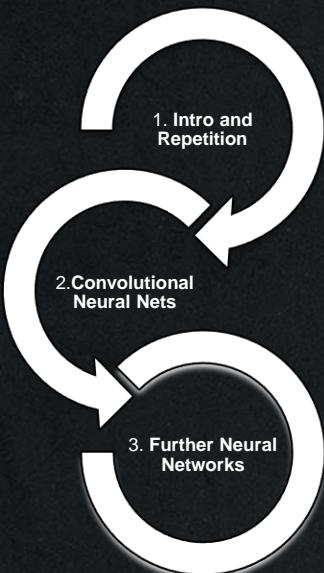
- [19] Patrick Ferdinand Christ, Mohamed Ezzeldin A Elshaer, Florian Ettlinger, et al. "Automatic liver and lesion segmentation in CT using cascaded fully convolutional neural networks and 3D conditional random fields". In: International Conference on Medical Image Computing and Computer-Assisted Springer. 2016, pp. 415–423.
- [20] Vincent Christlein, David Bernecker, Florian Höning, et al. "Writer Identification Using GMM Supervectors and Exemplar-SVMs". In: Pattern Recognition 63 (2017), pp. 258–267.
- [21] Florin Cristian Ghesu, Bogdan Georgescu, Tommaso Mansi, et al. "An Artificial Agent for Anatomical Landmark Detection in Medical Images". In: Medical Image Computing and Computer-Assisted Intervention — MICCAI 2016. Athens, 2016, pp. 229–237.
- [22] Jia Deng, Wei Dong, Richard Socher, et al. "Imagenet: A large-scale hierarchical image database". In: Computer Vision and Pattern Recognition, 2009. CVPR 2009. IEEE Conference IEEE. 2009, pp. 248–255.
- [23] A. Karpathy and L. Fei-Fei. "Deep Visual-Semantic Alignments for Generating Image Descriptions". In: ArXiv e-prints (Dec. 2014). arXiv: 1412.2306 [cs.CV].

References



- [24] Y LeCun, L Bottou, Y Bengio, et al. "Gradient-based Learning Applied to Document Recognition". In: Proceedings of the IEEE 86.11 (Nov. 1998), pp. 2278–2324. arXiv: 1102.0183.
- [25] Min Lin, Qiang Chen, and Shuicheng Yan. "Network in network". In: International Conference on Learning Representations. Banff, Canada, Apr. 2014. arXiv: 1102.0183.
- [26] Olga Russakovsky, Jia Deng, Hao Su, et al. "ImageNet Large Scale Visual Recognition Challenge". In: International Journal of Computer Vision 115.3 (Dec. 2015), pp. 211–252.
- [27] Karen Simonyan and Andrew Zisserman. "Very Deep Convolutional Networks for Large-Scale Image Recognition". In: International Conference on Learning Representations (ICLR). San Diego, May 2015. arXiv: 1409.1556.
- [28] Rupesh Kumar Srivastava, Klaus Greff, Urgen Schmidhuber, et al. "Training Very Deep Networks". In: Advances in Neural Information Processing Systems 28. Curran Associates, Inc., 2015, pp. 2377–2385. arXiv: 1507.06228.
- [29] C. Szegedy, Wei Liu, Yangqing Jia, et al. "Going deeper with convolutions". In: 2015 IEEE Conference on Computer Vision and Pattern Recognition (CVPR). June 2015, pp. 1–9.
- [30] C. Szegedy, V. Vanhoucke, S. Ioffe, et al. "Rethinking the Inception Architecture for Computer Vision". In: 2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR). June 2016, pp. 2818–2826.

References



- [31] Christian Szegedy, Sergey Ioffe, and Vincent Vanhoucke. "Inception-v4, Inception-ResNet and the Impact of Residual Connections on Learning". In: Thirty-First AAAI Conference on Artificial Intelligence (AAAI-17) Inception-v4, San Francisco, Feb. 2017. arXiv: 1602.07261.
- [32] Andreas Veit, Michael J Wilber, and Serge Belongie. "Residual Networks Behave Like Ensembles of Relatively Shallow Networks". In: Advances in Neural Information Processing Systems 29. Curran Associates, Inc., 2016, pp. 550–558. A.
- [33] Sergey Zagoruyko and Nikos Komodakis. "Wide Residual Networks". In: Proceedings of the British Machine Vision Conference (BMVC). BMVA Press, Sept. 2016, pp. 87.1–87.12.
- [34] K Zhang, M Sun, X Han, et al. "Residual Networks of Residual Networks: Multilevel Residual Networks". In: IEEE Transactions on Circuits and Systems for Video Technology PP.99 (2017), p. 1.

researchgate



Further online lectures

<http://cs229.stanford.edu>

Prof. Andrew Ng

Prof. Andreas Maier, FAU Erlangen:

https://www.youtube.com/watch?v=WJ-Es3gtR-M&list=PLpOGQvPCDQzvgpD3S0vTy7bJe2pf_yJFj&index=26